




Genome assembly

Meleshko Dmitry

meleshko.dmitrii@gmail.com

De novo genome assembly




De novo whole genome assembly




stack of NY Times, June 27, 2000



stack of NY Times, June 27, 2000
on a pile of dynamite




this is just hypothetical



so, what did the June 27, 2000 NY
Times say?

Whole genome assembly



Shortest common supersequence

Given a set of strings $\{s_1, \dots, s_n\}$, find a shortest string **S** containing each s_i as a substring

Shortest common supersequence

Given a set of strings $\{s_1, \dots, s_n\}$, find a shortest string **S** containing each s_i as a substring

Is NP-complete

Has nothing to do with real genome assembly problem

Why to assemble?

- NGS

- Billions of short reads
- Sequencing errors
- Contaminants

Hard to perform analysis

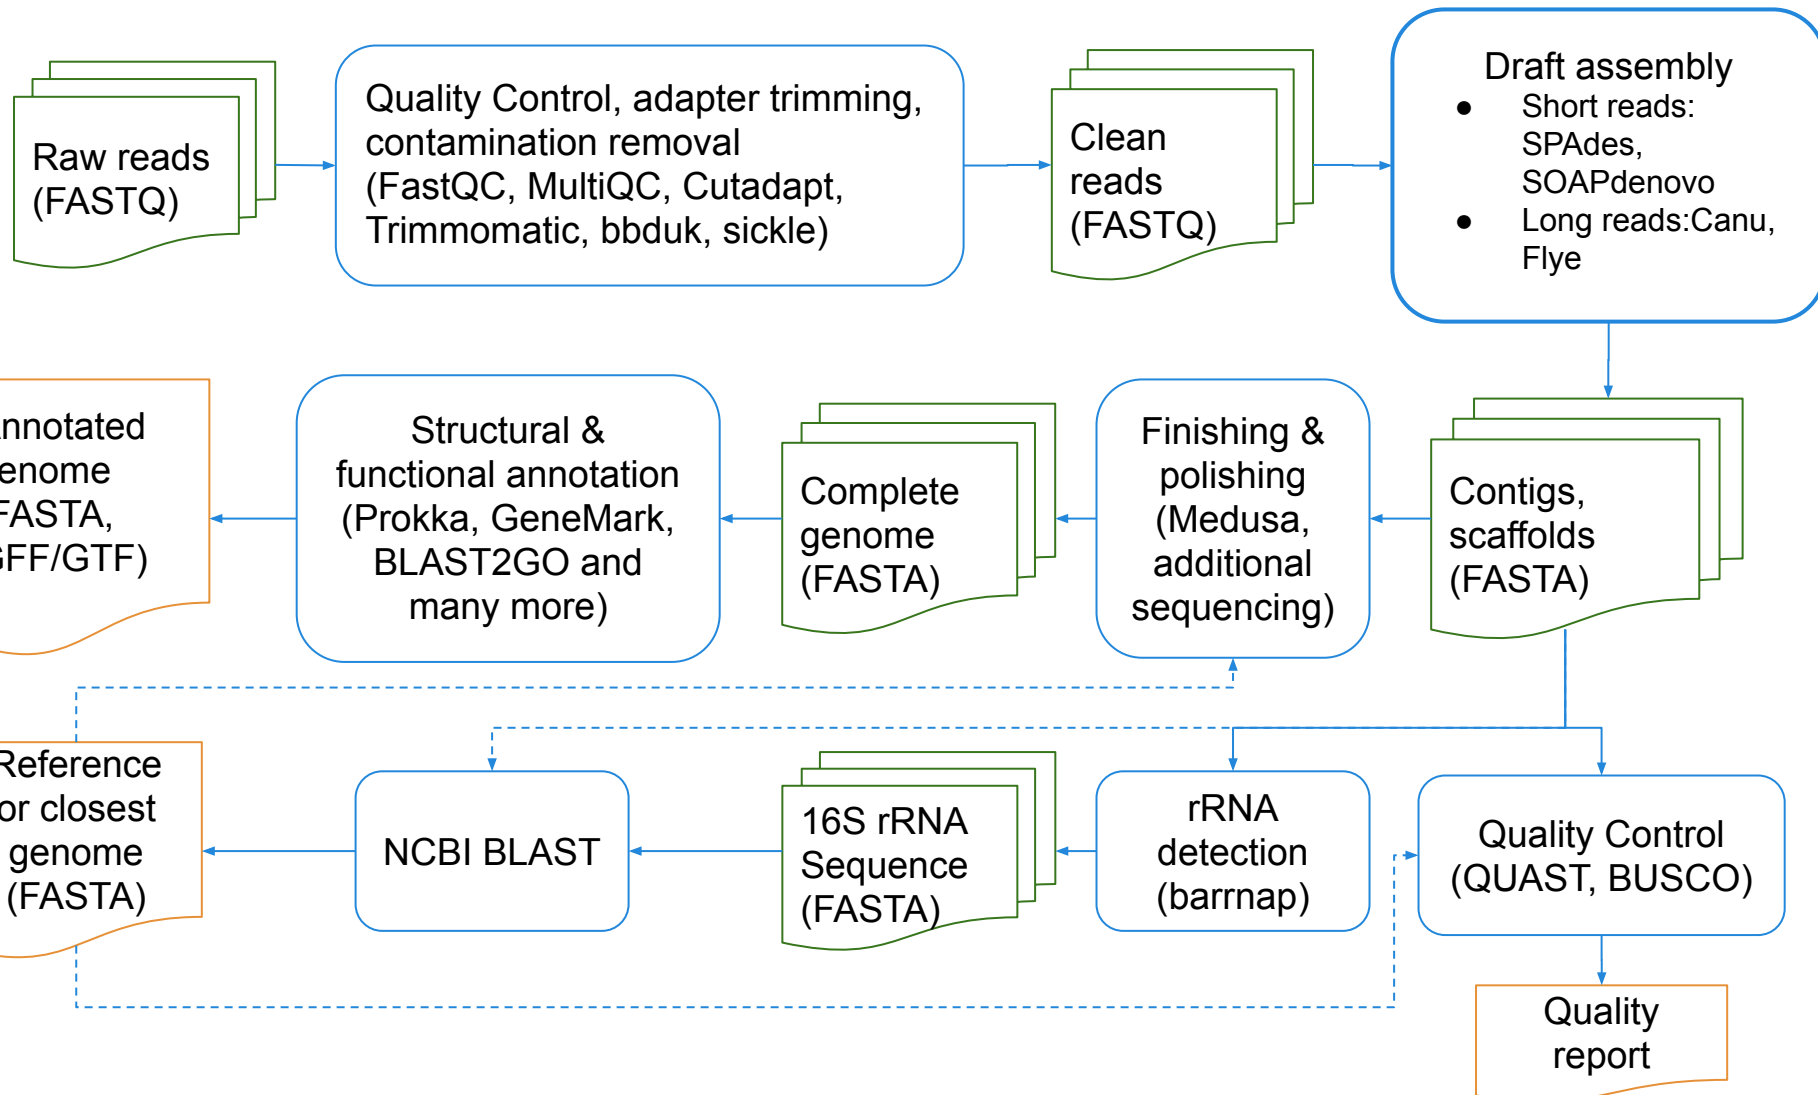
- Assembly

- ✓ Corrects sequencing errors
- ✓ Much longer sequences
- ✓ Each genomic region is presented only once
- ✗ May introduce errors

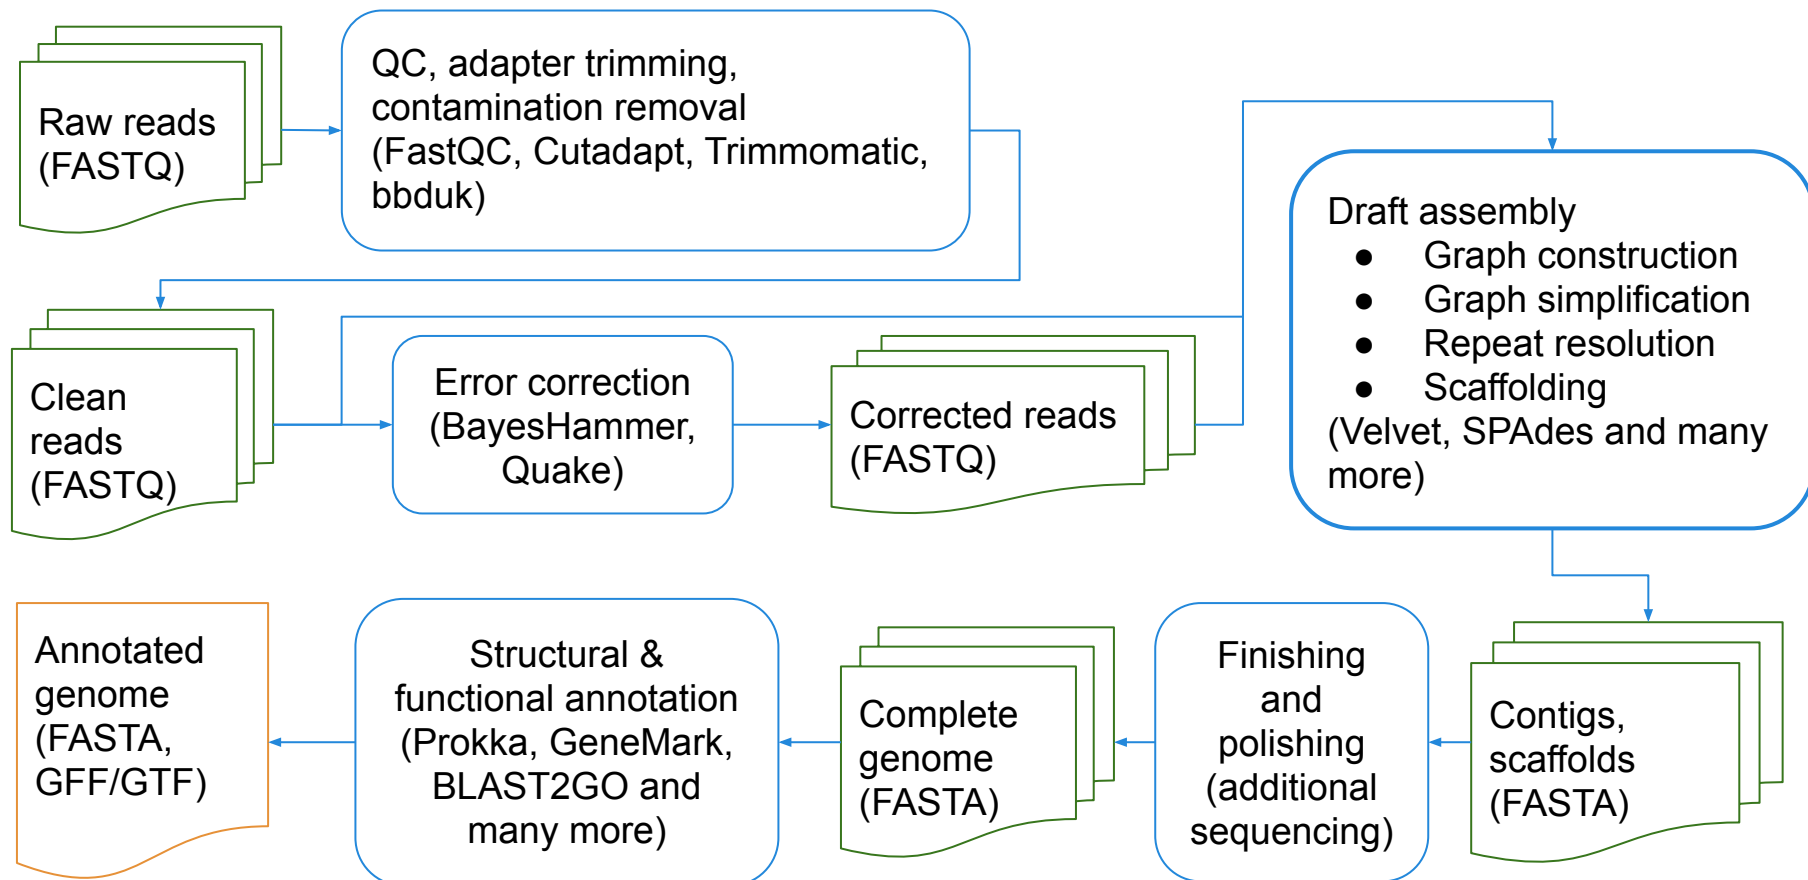
Assembly types

- *De novo* genome assembly
 - Long reads
 - Short reads
 - Hybrid
- Reference-assisted genome assembly
 - Closely related species
- Transcriptome assembly
 - *De novo*
 - Reference based

De novo genome assembly



De novo genome assembly



Assembling Sanger reads

Early days

- Sanger sequencing
 - Long reads
 - Low coverage
- Overlap-Layout-Consensus (OLC)
 - Find overlaps between all reads
 - Order reads
 - Merge into consensus sequence

Finding overlaps

- Align reads all-to-all
 - BLAST and similar algorithms
- Ignore "insufficient" overlaps


Finding overlaps

- Align reads all-to-all
 - BLAST and similar algorithms
- Ignore "insufficient" overlaps
 - At least 40bp
 - >94% similarity


Assembly example





Assembly example




Assembly example




Overlap graph




Overlap graph




Layout




Layout





Layout




Layout



Layout



Consensus



NGS and OLC

- Overlap-Layout-Consensus is not applicable
 - Hard to find overlaps between short reads
 - Impossible to scale to such amount of reads
- De Bruijn graph approach
(Pevzner et al., 2001)
- String Graph approach
(Meyers, 2005)



NGS era

De Bruijn graph in a nutshell

He that mischief hatches, mischief catches

I


Sequencing




**He that mischief
mischief hatches,
hatches, mischief
, mischief catches**

De Bruijn graph in a nutshell

, mischief catches
mischief hatches,
He that mischief
hatches, mischief



De Bruijn graph in a nutshell




De Bruijn graph

ACGTCCGTAA

De Bruijn graph

ACGTCCGTAA



$k=2$



De Bruijn graph

ACGTCGTA




$k=2$



De Bruijn graph

ACGTCCGTAA

$k=2$



De Bruijn graph

ACGTCGTA

$k=2$

AC

CG

GT

TC

De Bruijn graph

ACGTCCGTAA

$k=2$

AC

CG

GT

CC

TC

De Bruijn graph

ACGTC**CG**TAA

$k=2$

AC

CG

GT

CC

TC

De Bruijn graph

ACGTCCGTAA

$k=2$

AC

CG

GT

CC

TC

De Bruijn graph

ACGTCCGTAA

$k=2$

AC

CG

GT

TA

CC

TC

De Bruijn graph

ACGTCCGTAA

$k=2$

AC

CG

GT

TA

CC


TC

AA

De Bruijn graph

ACGTCCGTAA


k=2



De Bruijn graph

ACGTCCGTAA

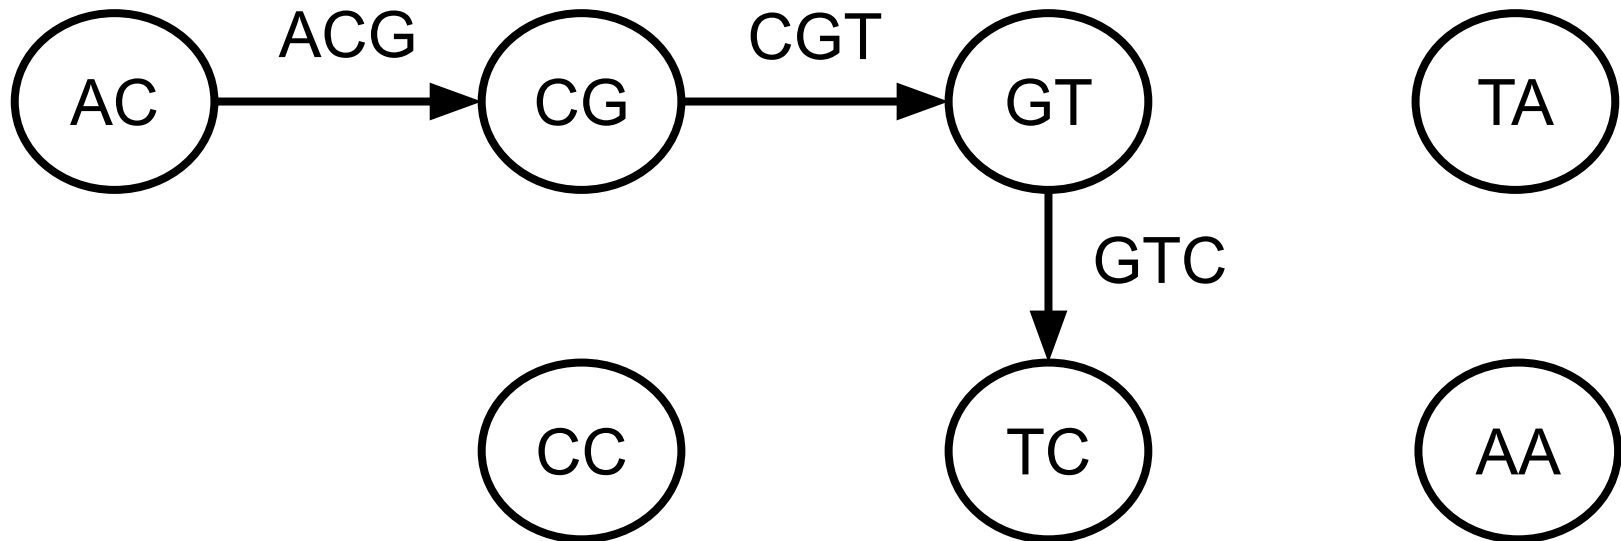
$k=2$



De Bruijn graph

ACGTCGTAA

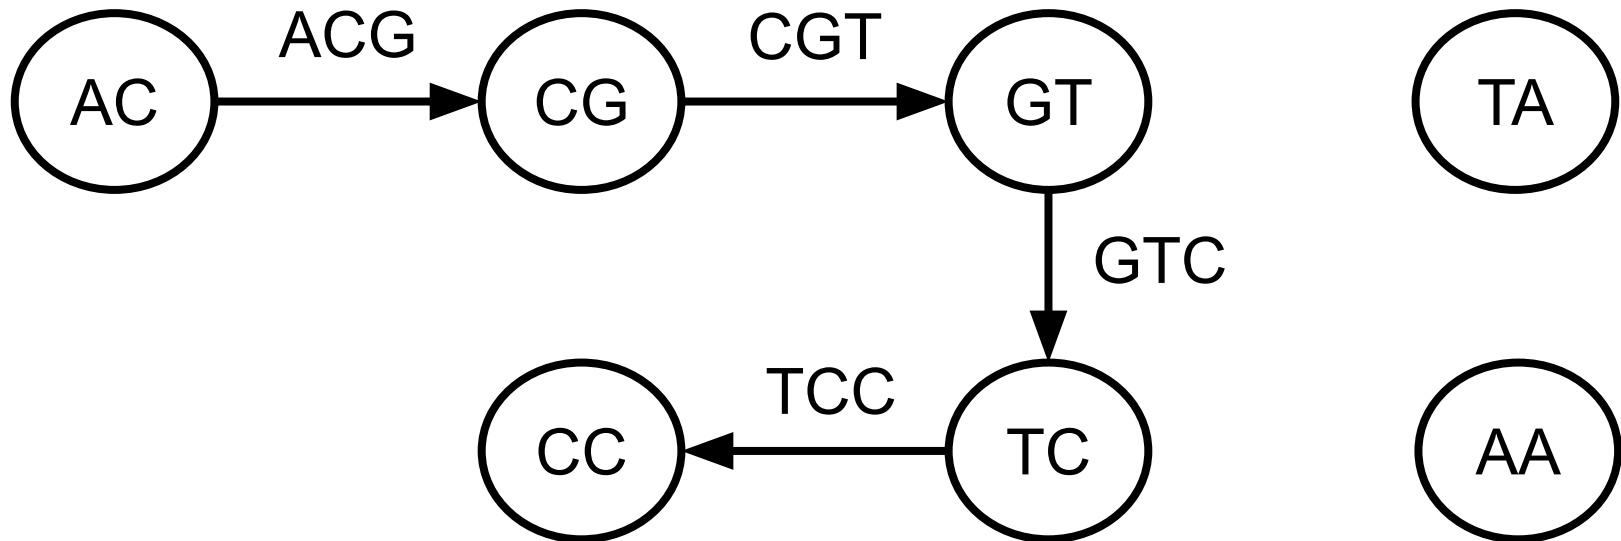
$k=2$



De Bruijn graph

ACGTCCGTAA

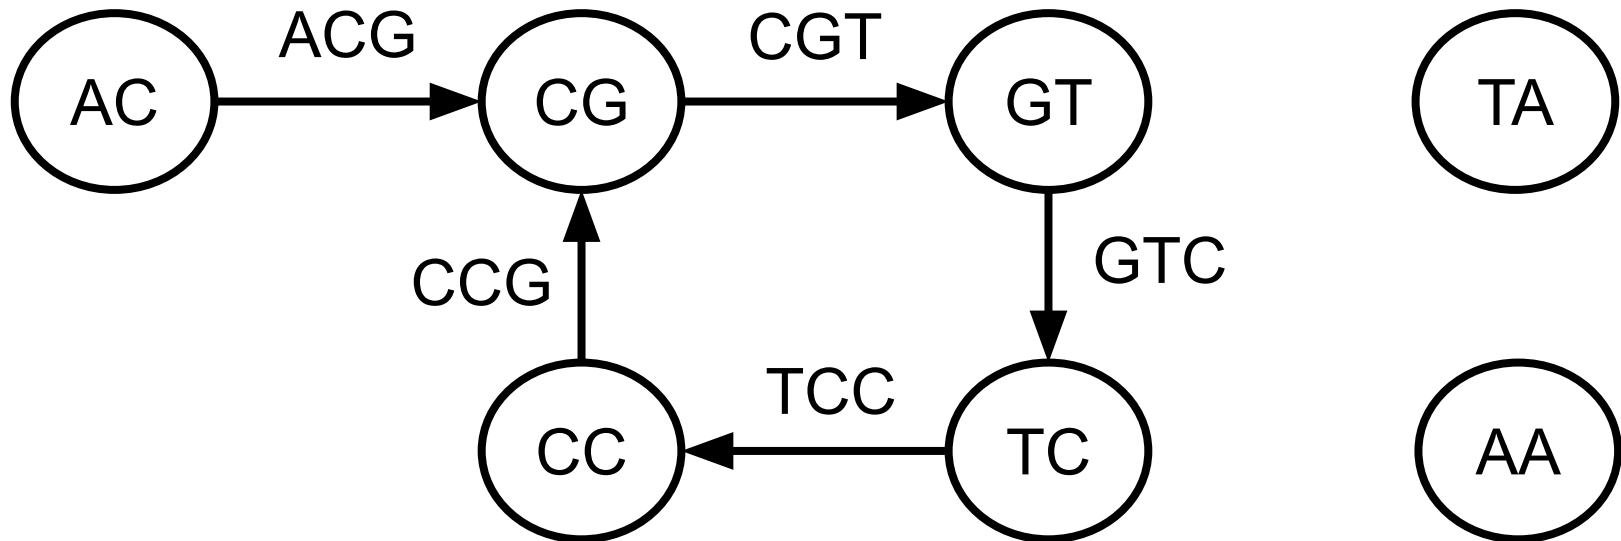
$k=2$



De Bruijn graph

ACGTCCGTAA

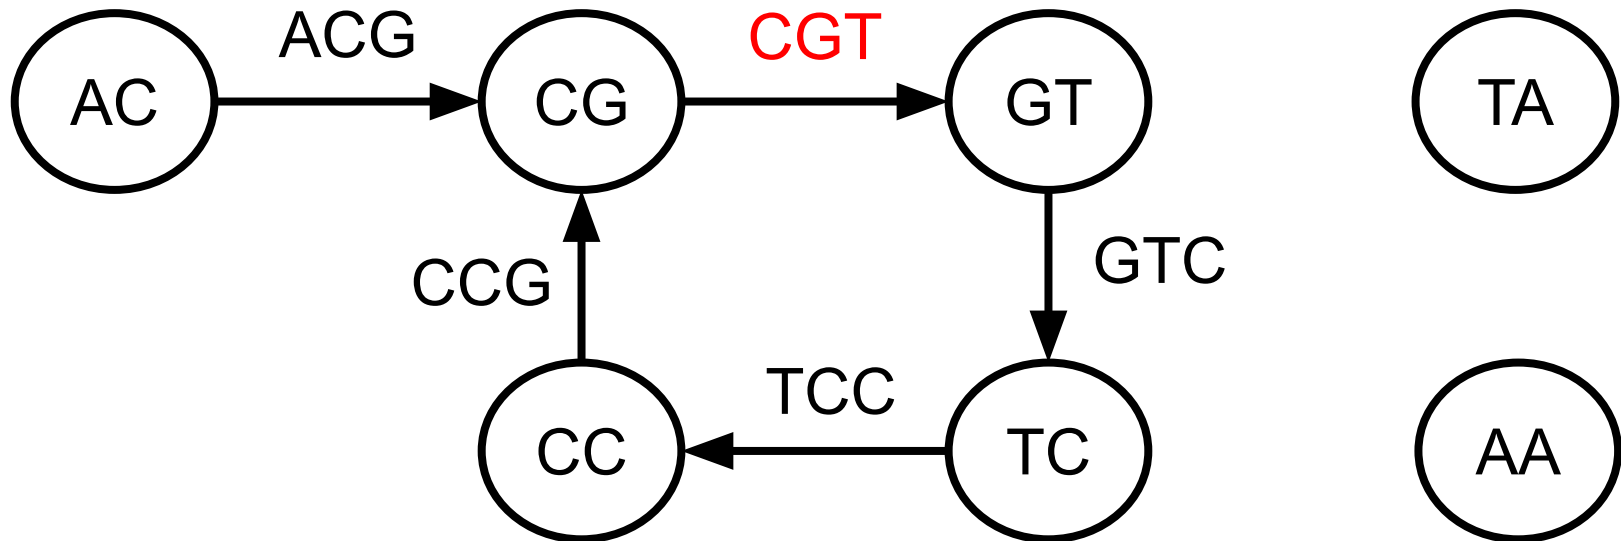
$k=2$



De Bruijn graph

ACGTC**CGT****A**A

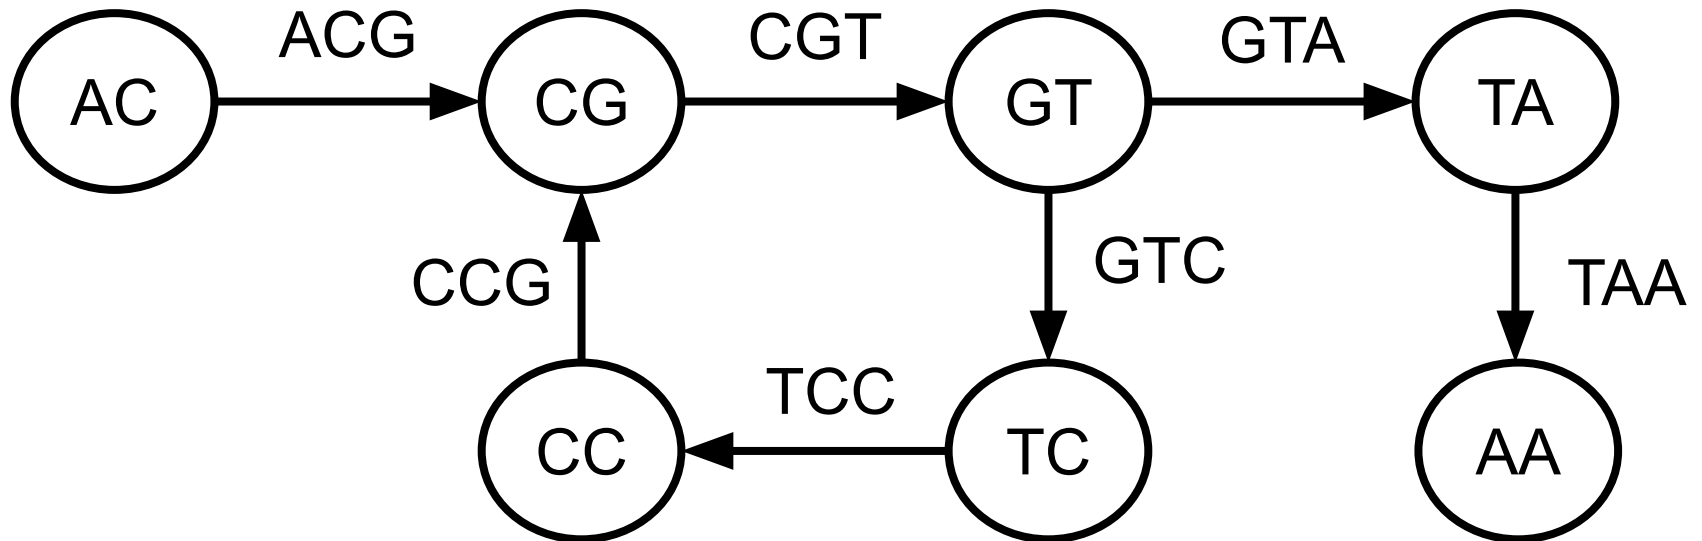
k=2



De Bruijn graph

ACGTCCGTAA

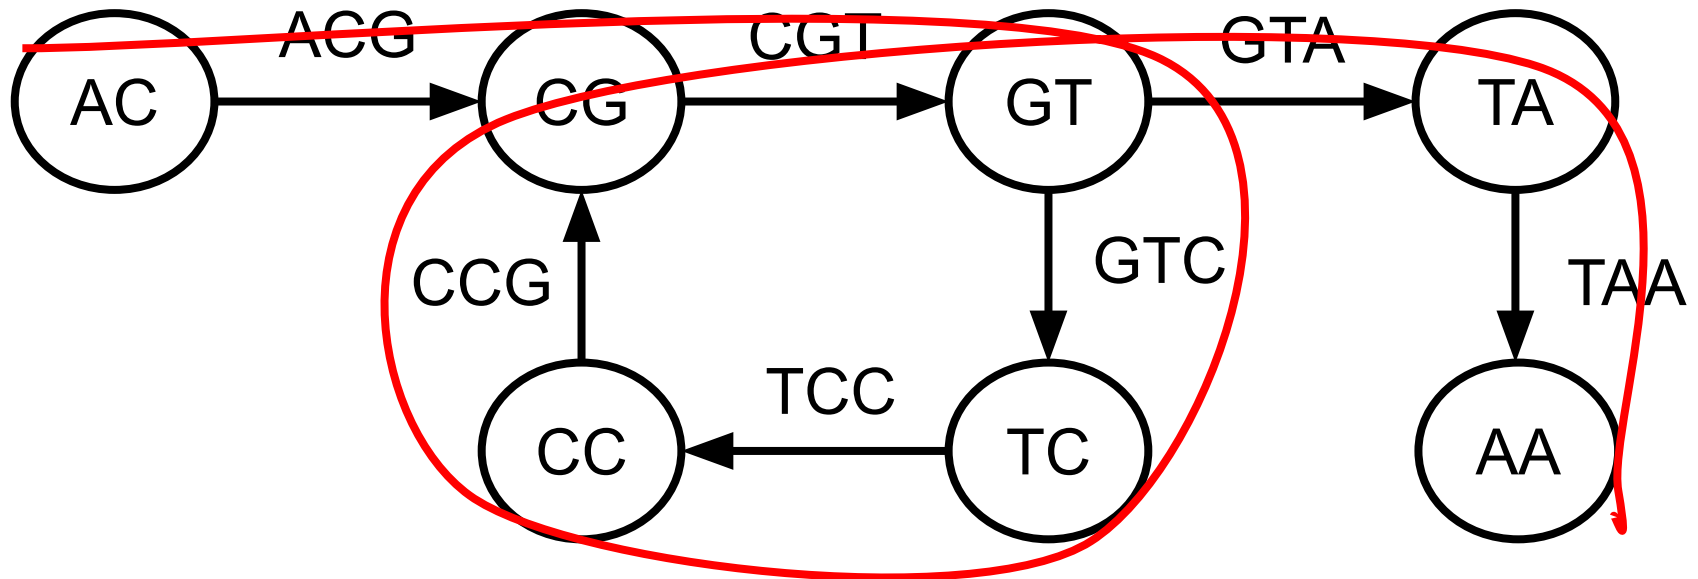
$k=2$



De Bruijn graph

ACGTCCGTAA


$k=2$



Condensed de Bruijn graph

ACGTCCGTAA

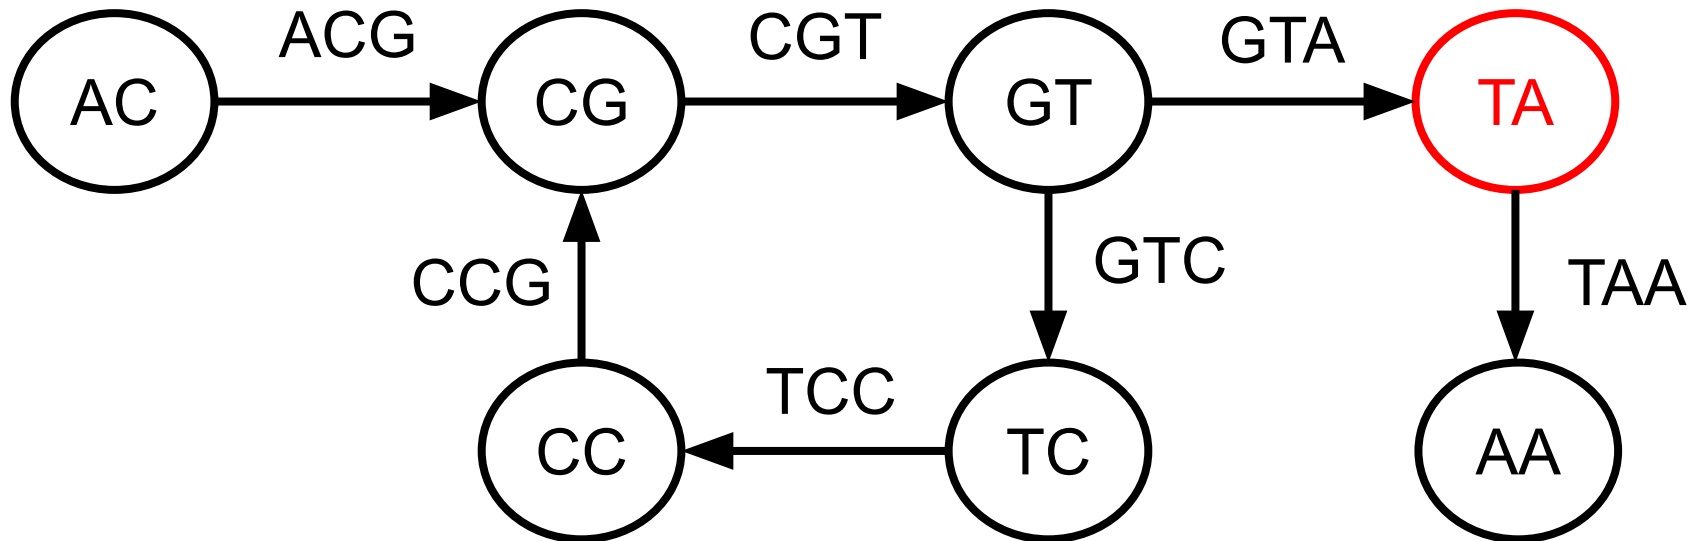
k=2



Condensed de Bruijn graph

ACGTCCGTAA

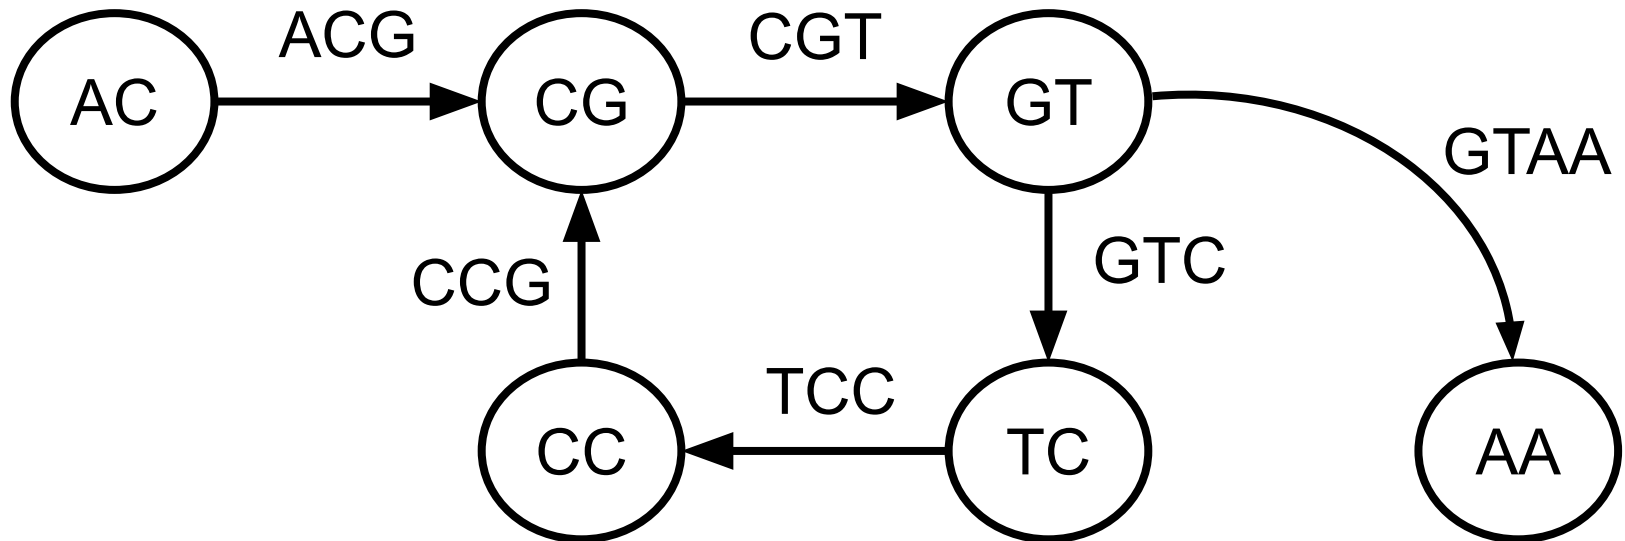
$k=2$



Condensed de Bruijn graph

ACGTCCGTAA

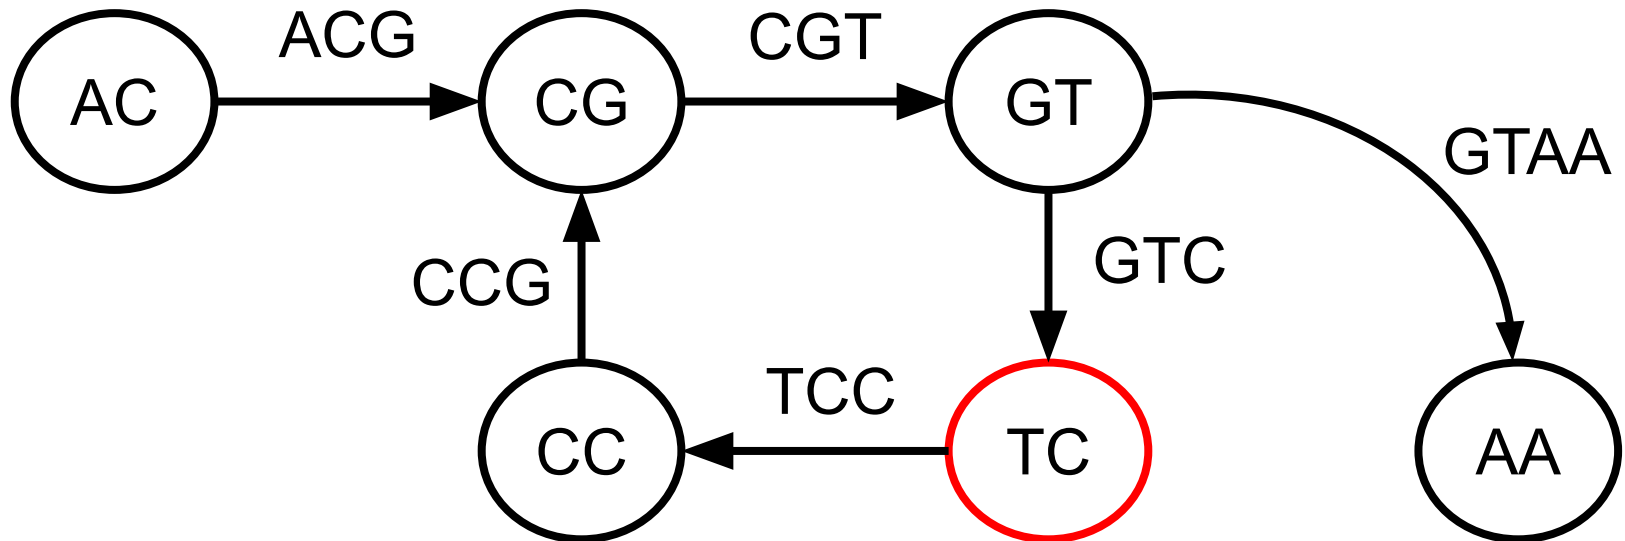
$k=2$



Condensed de Bruijn graph

ACGTCCGTAA

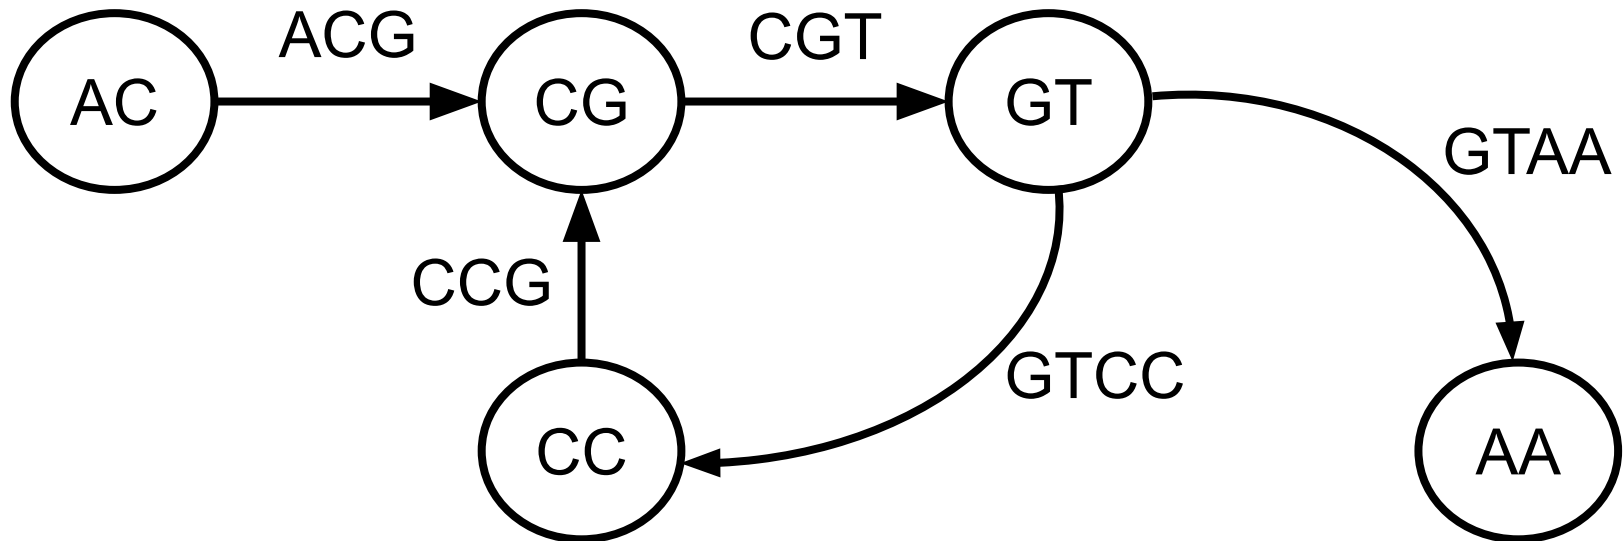
$k=2$



Condensed de Bruijn graph

ACGTCCGTAA

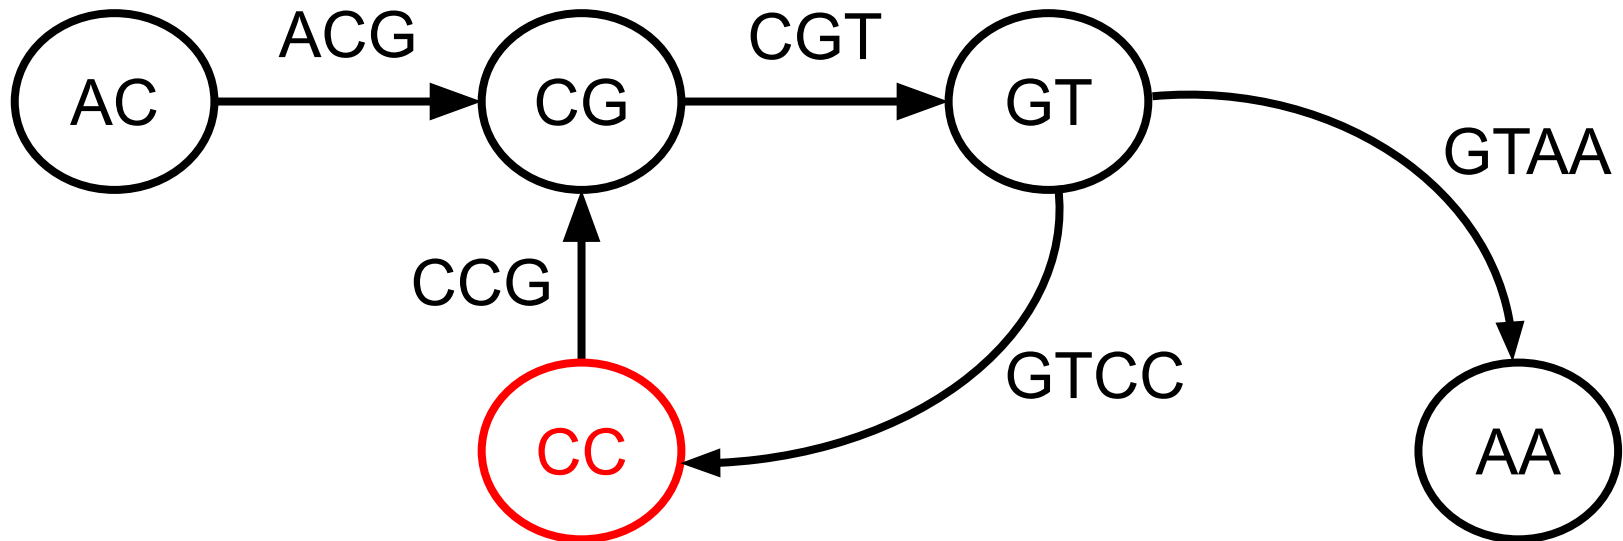
k=2



Condensed de Bruijn graph

ACGTCCGTAA


$k=2$



Condensed de Bruijn graph

ACGTCCGTAA


$k=2$



Condensed de Bruijn graph

ACGTCCGTAA


$k=2$



Repeats in de Bruijn graph

ACGTCCGTAA


$k=2$



Repeats in de Bruijn graph

ACGTCCGTAA


k=2



Repeats in de Bruijn graph

ACGTCCGTAA


$k=2$



Eulerian path with multiplicities

ACGTCCGTAA

$k=2$



More examples

CCGTTG

TGCAGG

GTTGCA

k=3

More examples

CCG

CCGTTG

TGCAGG

GTTGCA

k=3

More examples

CCG

CGT

CCGTTG
TGCAGG
GTTGCA

k=3

More examples

CCG

CGT

GTT

CCGTTG
TGCAGG
GTTGCA

k=3

More examples

CCG

CGT

GTT

TTG

CCGTTG

TGCAGG

GTTGCA

k=3

More examples

CCG

CGT

CCGTTG

TGCAGG

GTTGCA

k=3

GTT

TTG

TGC

GCA

CAG

AGG

More examples

CCG

CGT

CCGTTG

CCGT

TGCAGG

GTT

TTG

GTTGCA

k=3

TGC

GCA

CAG

AGG

More examples

CCG

CGT

CCGTTG

TGCAGG

GTTGCA

k=3

GTT

TTG

TGC

GCA

CAG

AGG

CCGT

CGTT

More examples

CCG

CGT

CCGTTG

TGCAGG

GTTGCA

k=3

GTT

TTG

TGC

GCA

CAG

AGG

CCGT

CGTT

GTTG

What about real data?

CCG

CGT

CCGTTG
TGCAGG
GTTGCA

k=3

GTT

TTG

TGC

GCA

CAG

AGG

CCGT

CGTT

GTTG


TGCA

GCAG

CAGG

TTGC

More examples



GTT

TTG

TGC


GCA

CAG


AGG

CCGTTG
TGCAGG
GTTGCA

k=3




What about real data?




CCGTTG
TGCAGG
GTTGCA

k=3




More examples




CCGTTG
TGCAGG
GTTGCA

k=3




More examples




CCGTTG
TGCAGG
GTTGCA

k=3



More examples



CCGTTG
TGCAGG
GTTGCA

k=3



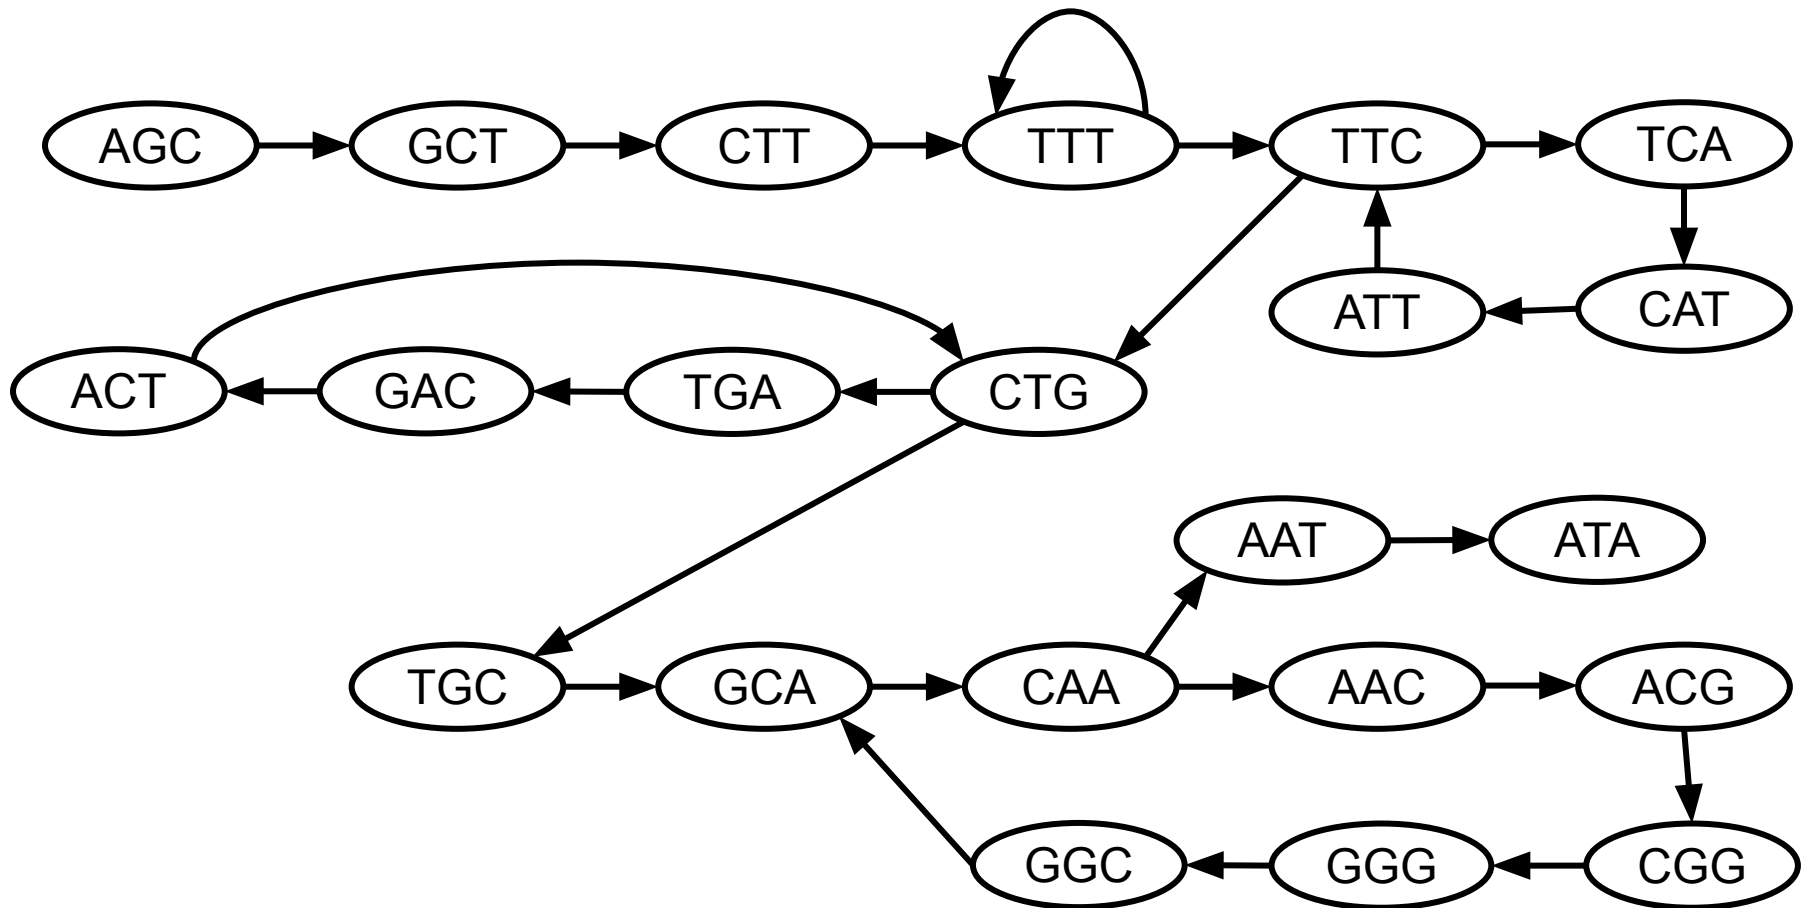
Does k-mer size matter?

One more example

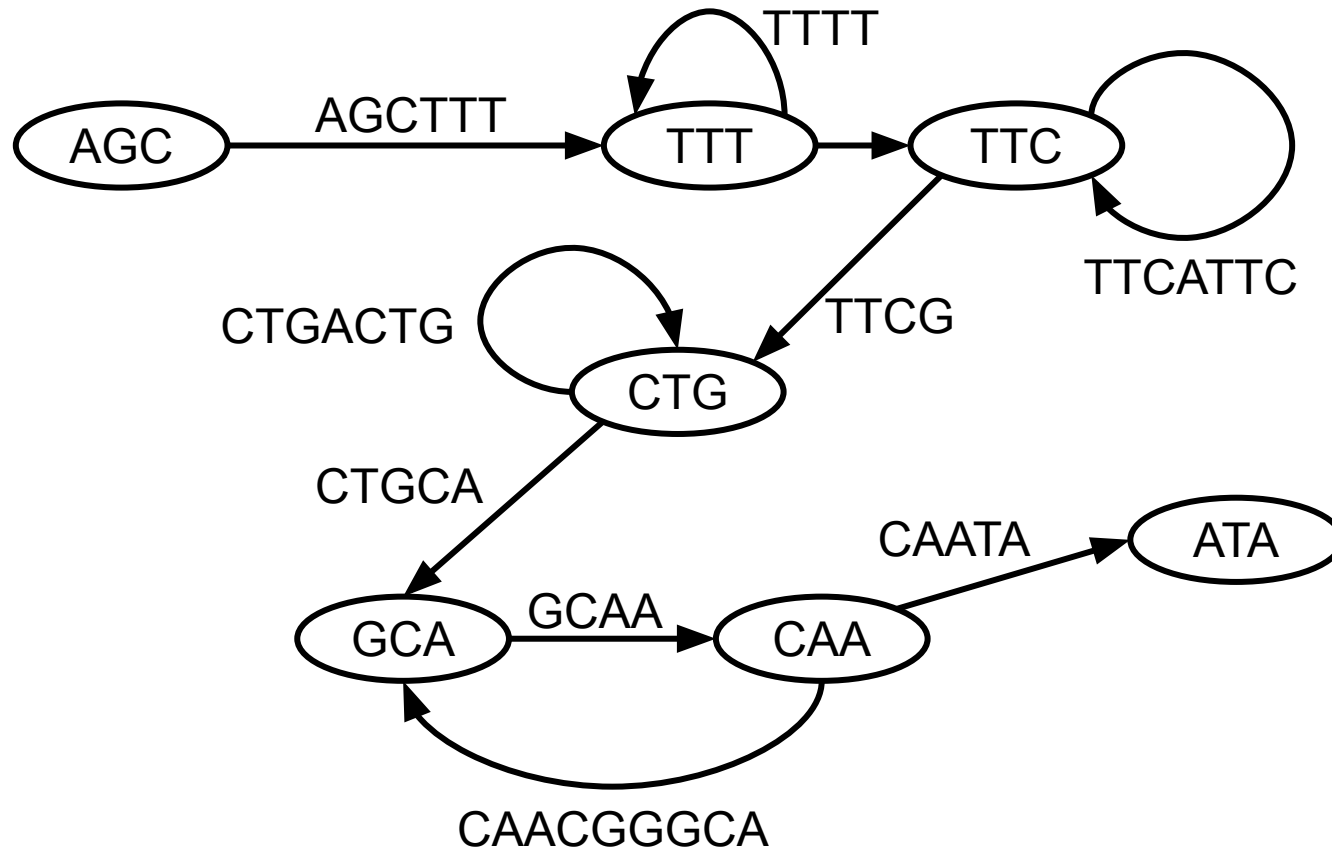
TTTCATTC	AACGGGCA
AGCTTTTC	CTGCAACG
GGGCAATA	TGACTGCA
CATTCTGA	

K = 3


One more example



One more example



One more example




AGCTTTTTCATTCTGACTGCAACGGGCAATA

One more example


TTTCATTC	AACGGGCA
AGCTTTTC	CTGCAACG
GGGCAATA	TGACTGCA
CATTCTGA	

K = 4

One more example



One more example



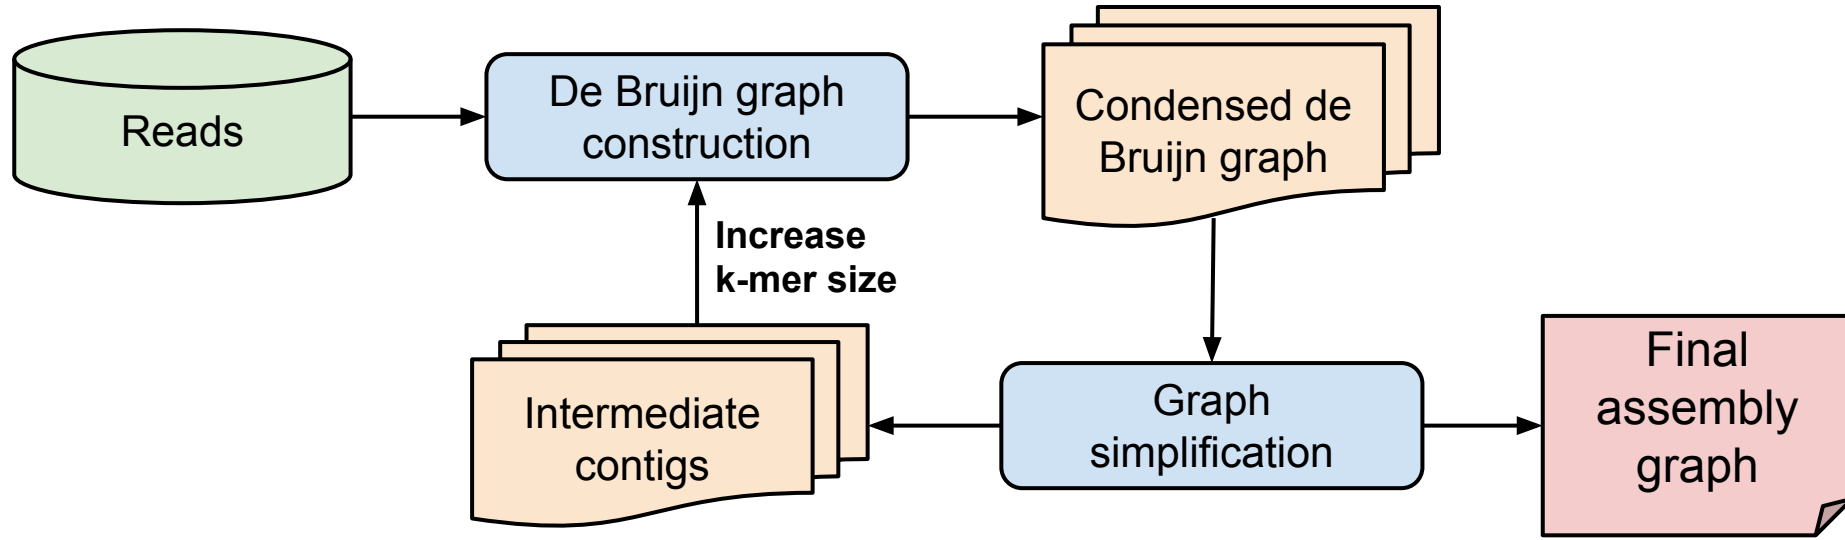
AGCTTTTTCATTCTGACTGCAACGGGGCAATA

How to select the k -mer size?

How to select the k -mer size?

- Small k
 - Complex graph
 - Hard to resolve repeats
- Large k
 - Gaps in the assembly
- For normal data sets $k = \textit{ReadLength} / 2 + \varepsilon$

Iterative SPAdes run



- Smaller k -mer sizes are needed for reconstructing low-coverage regions
- Larger k -mer sizes are needed for resolving short repeats

DNA is double-stranded

DNA is double-stranded

- Add k-mer and its reverse complement
- Use odd k to avoid self-complement vertices
 - $\text{rc}(\text{AATTT}) = \text{AAATT}$
 - $\text{rc}(\text{AATT}) = \text{AATT}$

Removing sequencing errors

Sequencing errors in de Bruijn graph

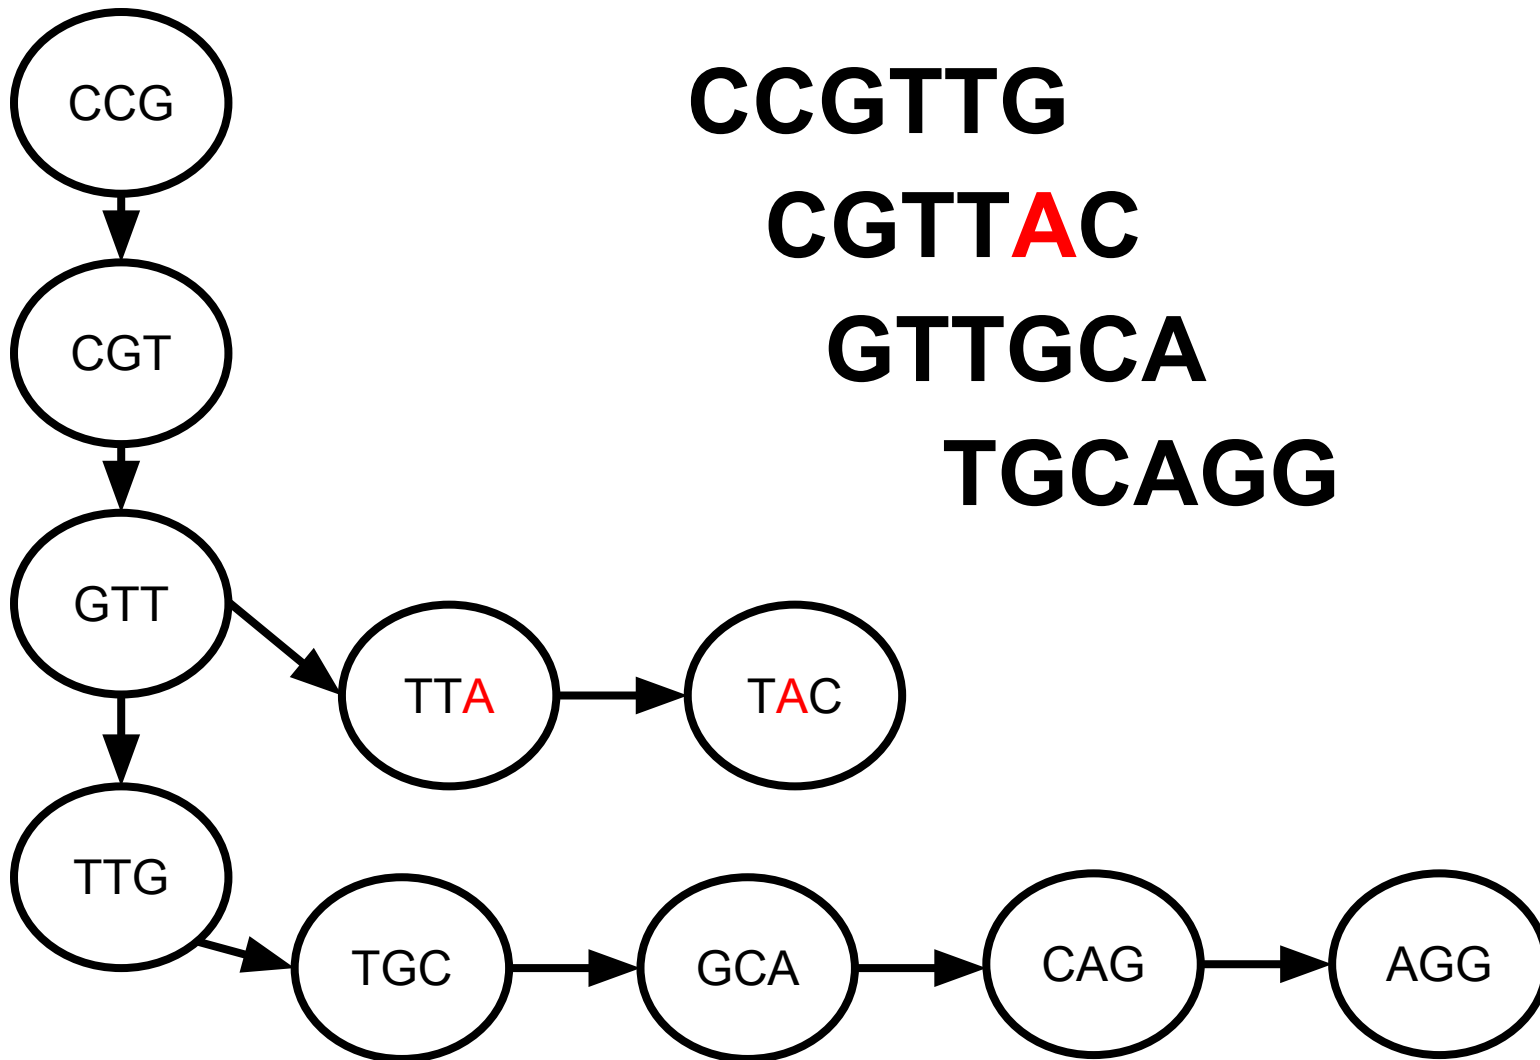
CCGTTG

CGTTAC

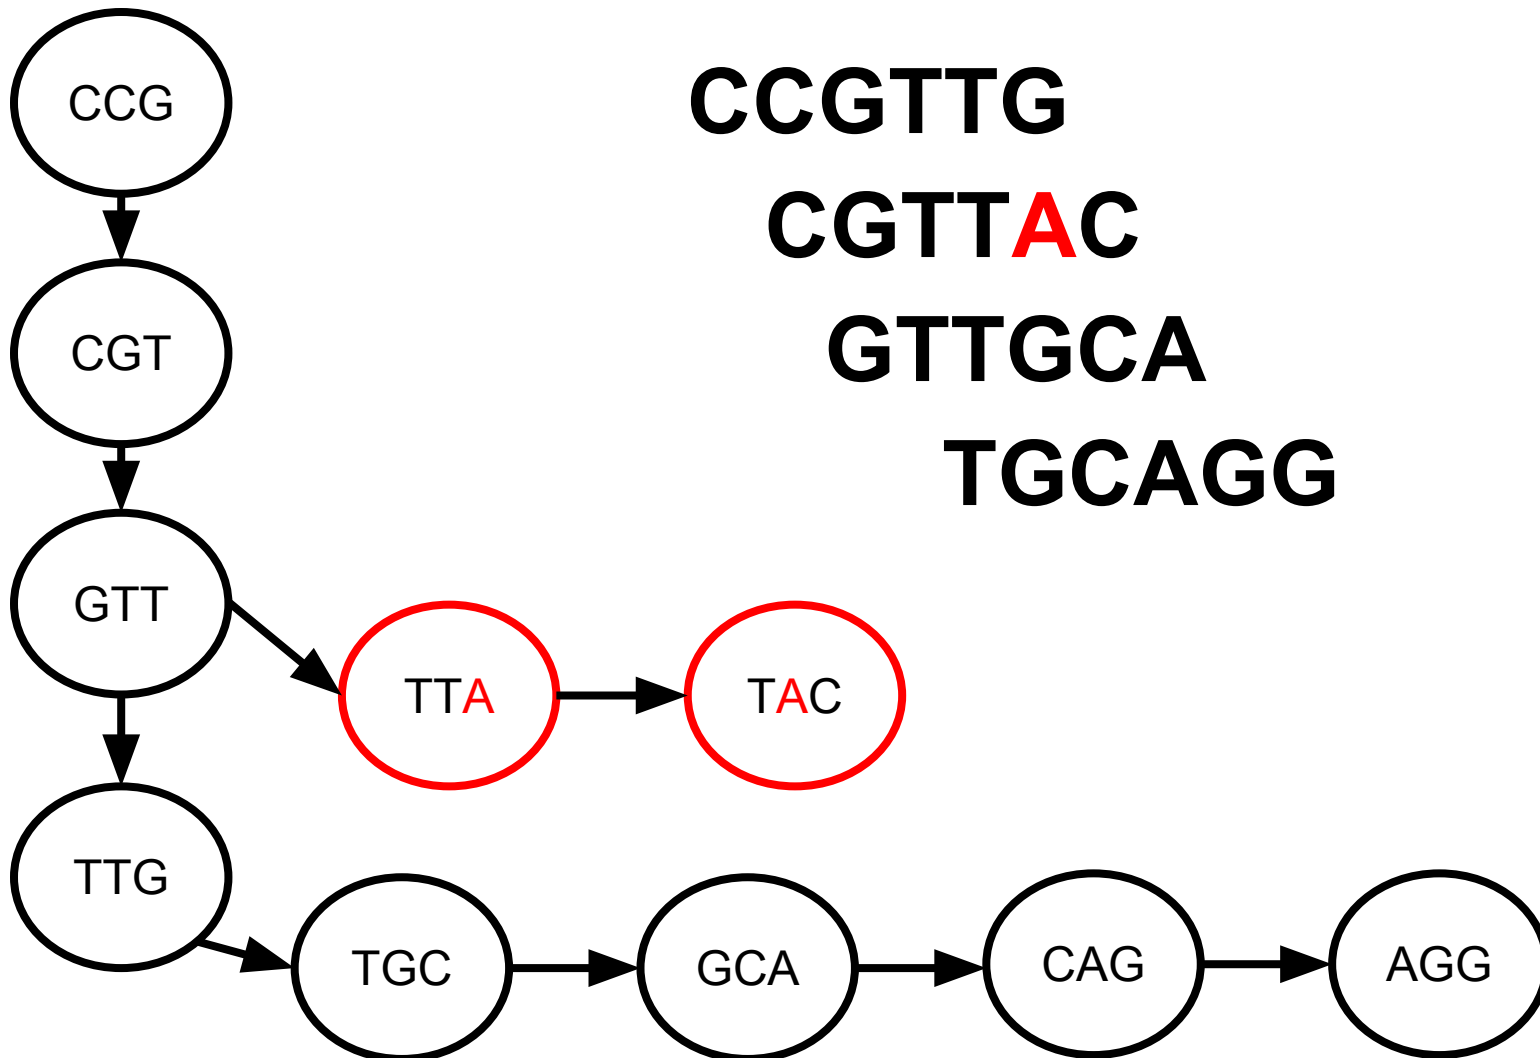
GTTGCA

TGCAGG

Sequencing errors in de Bruijn graph

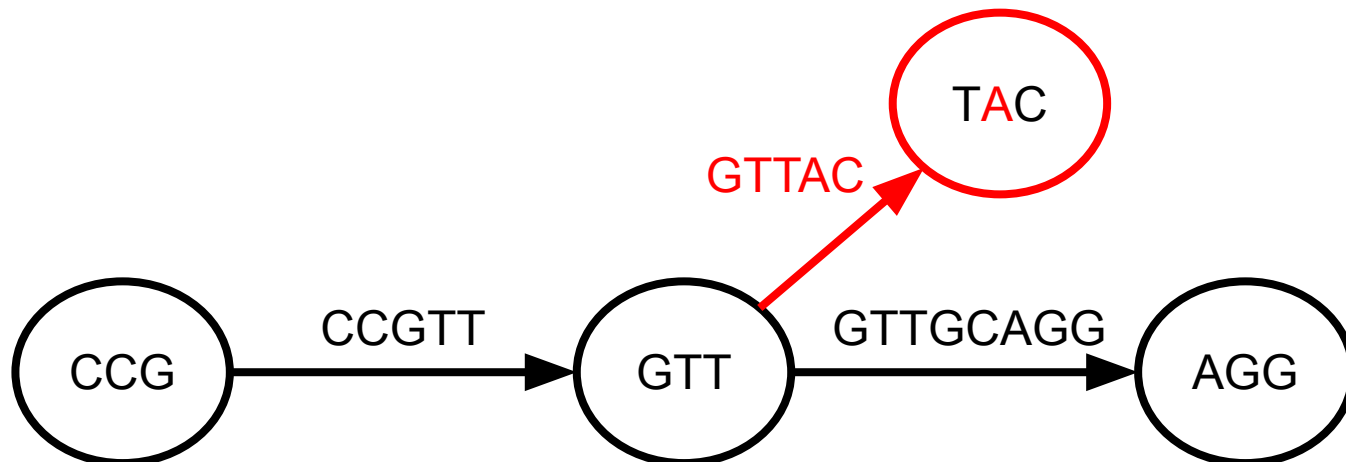


Sequencing errors in de Bruijn graph



Sequencing errors in de Bruijn graph

CCGTTG
CGTTAC
GTTGCA
TGCAGG

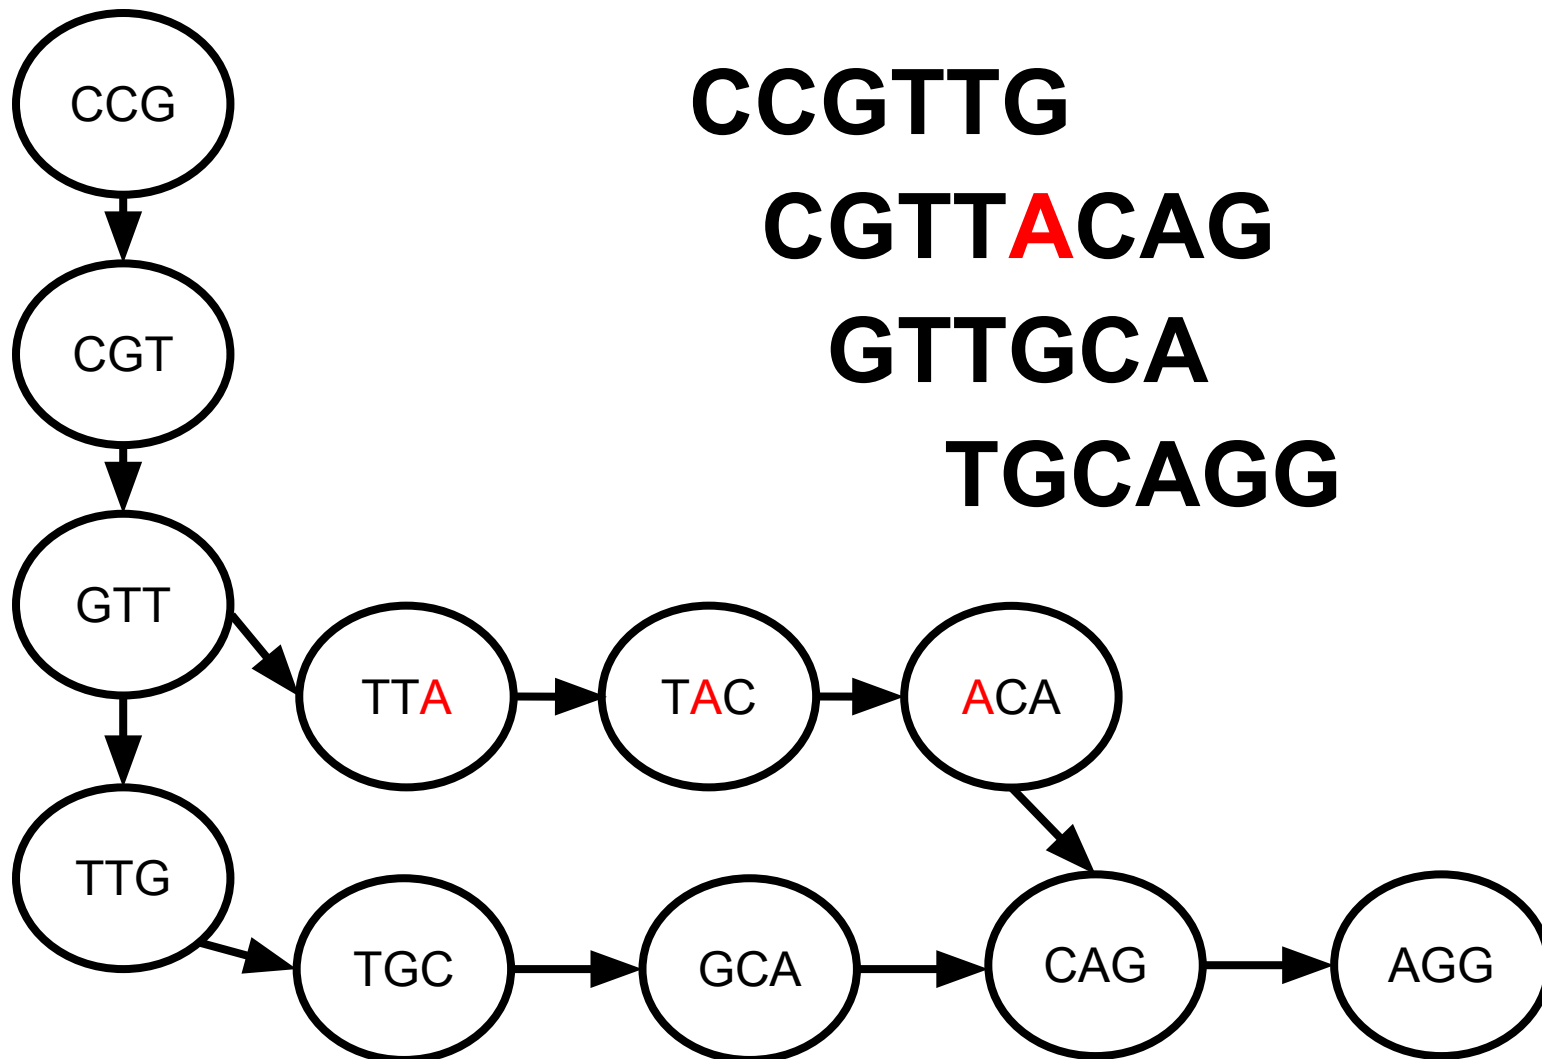


How to remove a tip?

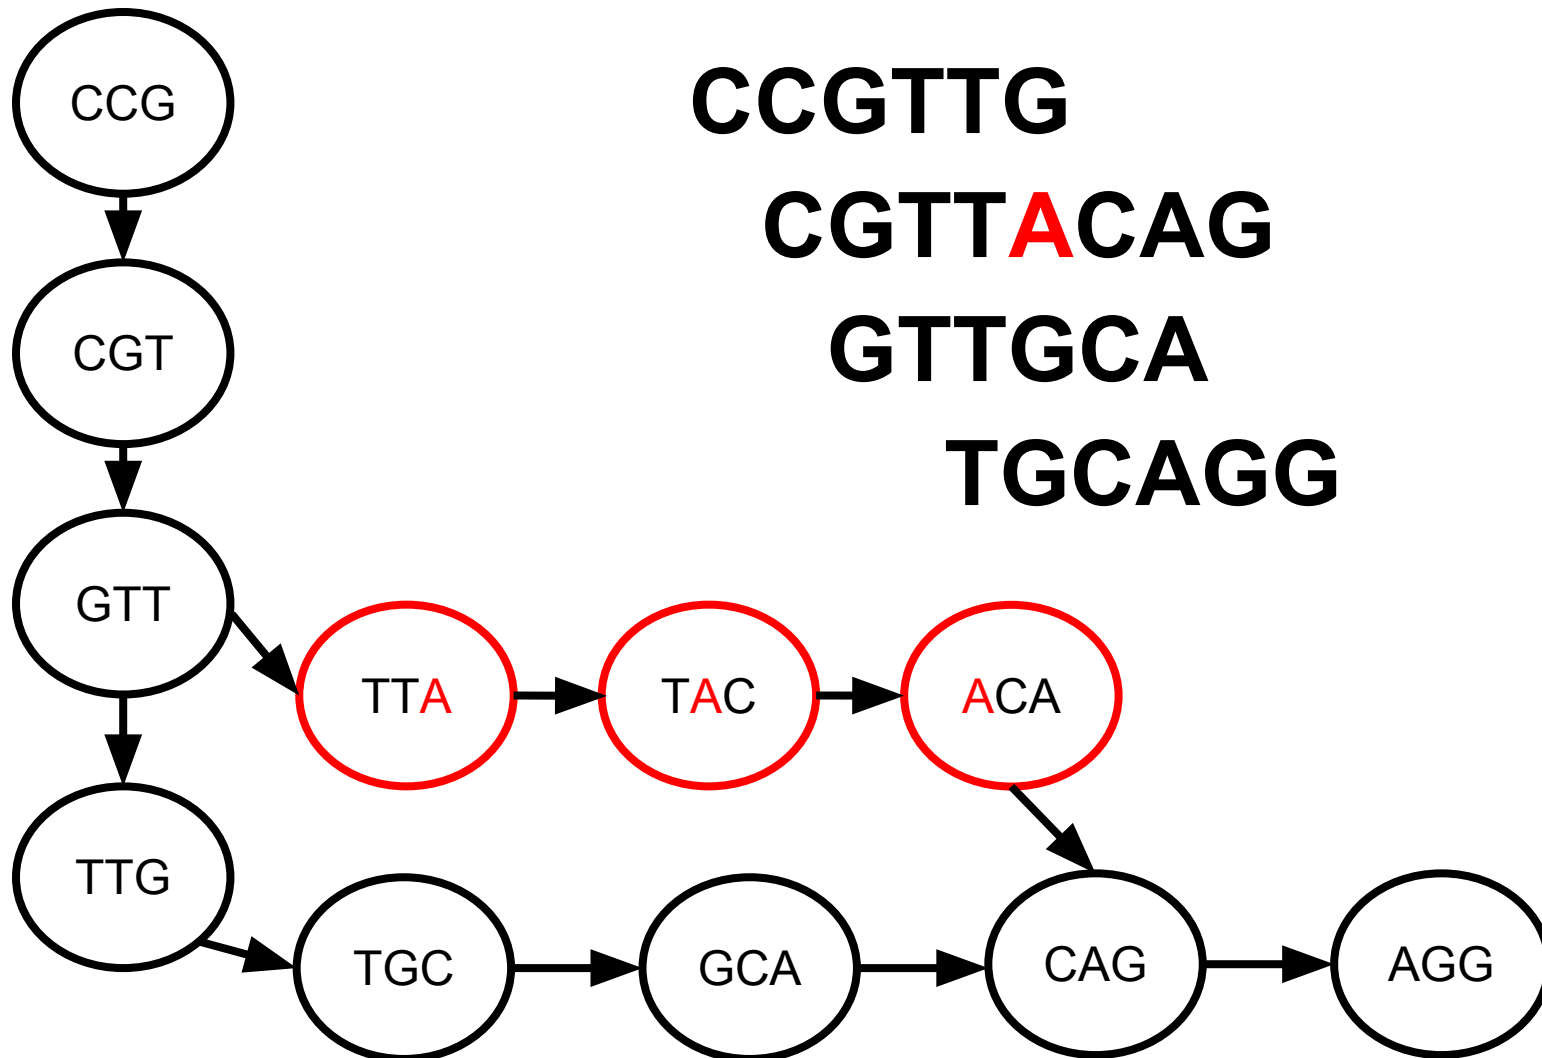
How to remove a tip?

- Short length (usually less than $2 * k$)
- Low coverage in respect to the main (correct path)
- Long length or high coverage — more likely to indicate a coverage gap

More about sequencing errors



More about sequencing errors



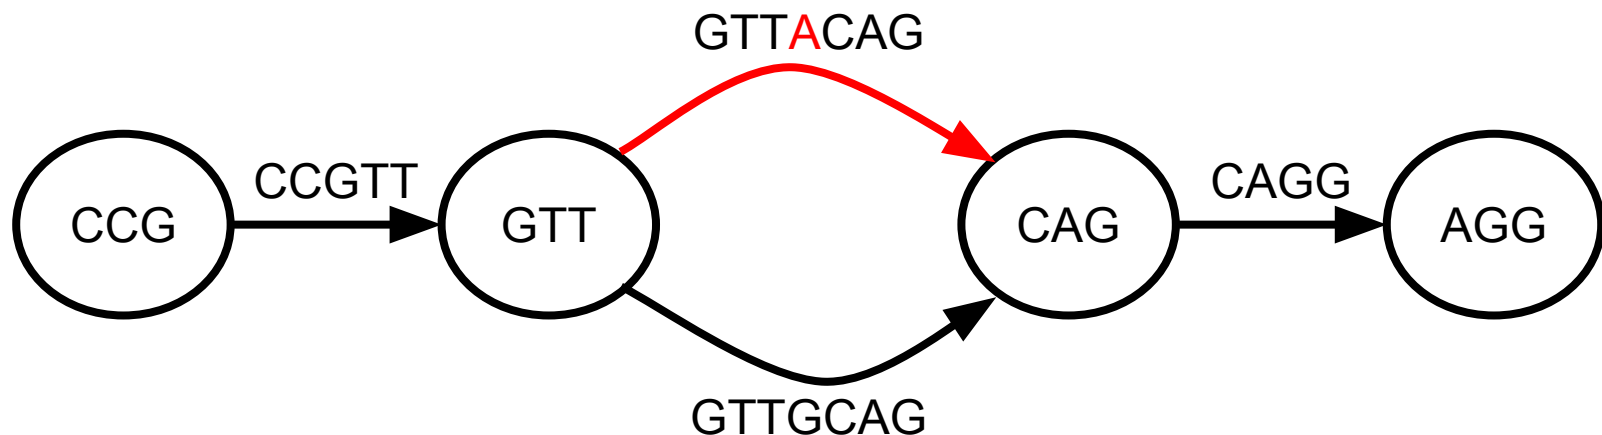
More about sequencing errors

CCGTTG

CGTTACAG

GTTGCA

TGCAGG




More about sequencing errors

CCGTTG

CGTTACAG

GTTGCA

TGCAGG

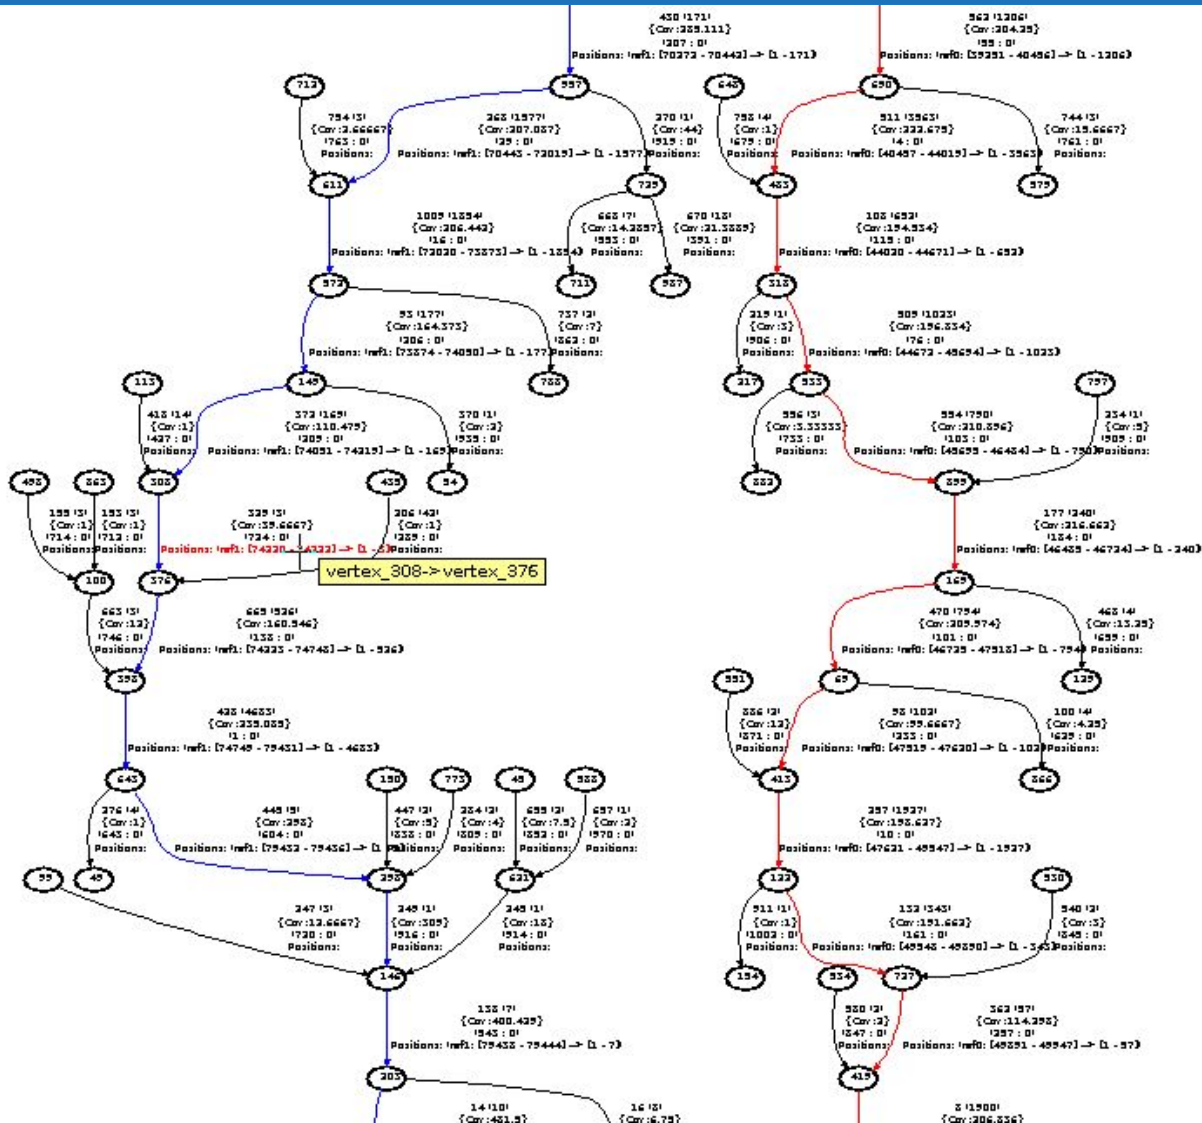


And what about bulges?

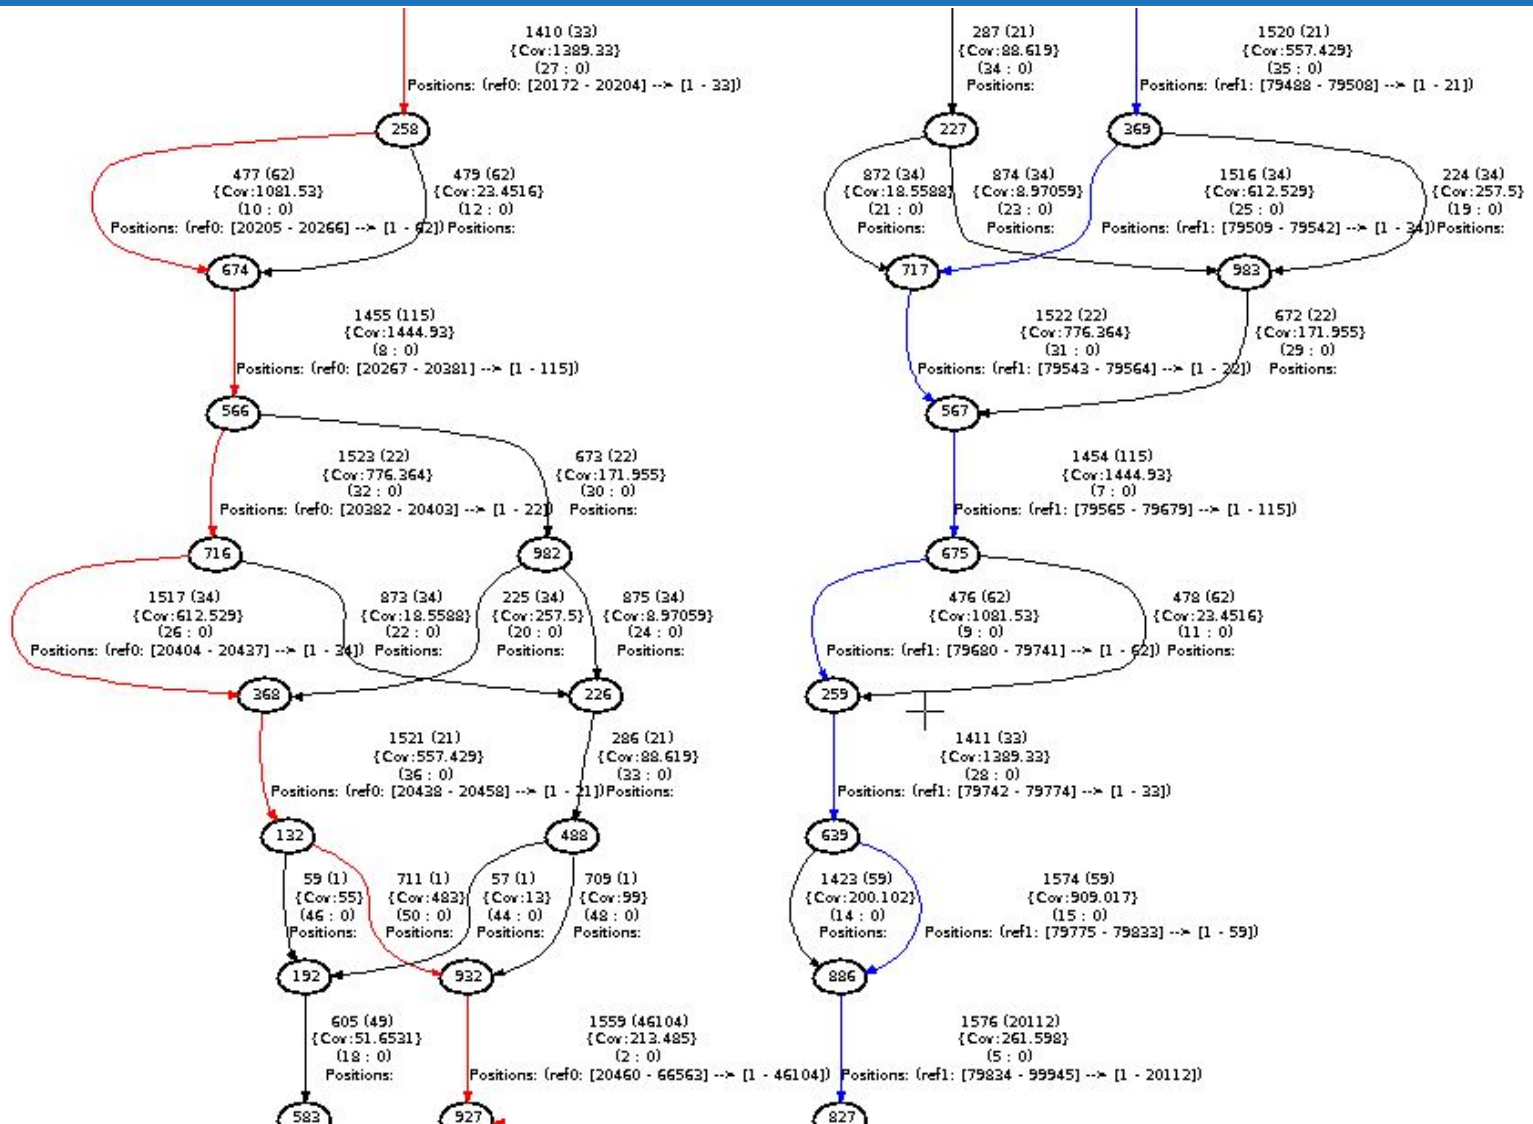
And what about bulges?

- Erroneous path has lower coverage than correct one
- Rather small length
- In case of similar coverage or bigger — more likely to be result of diploidy

Real life



Real life



Velvet assembler simplification

- Tip clipping
- Bulge removal
- Removing erroneous connections


Velvet tip clipping

- Remove only if shorter than $2 * k$
- Coverage is lower than of any alternative paths
- Iteratively process over the graph until no tips are left


Velvet “tour bus” algorithm

- Distance between vertices A and B is
$$D(A, B) = \textit{length}(E_{AB}) / \textit{coverage}(E_{AB})$$
 - Allows to go through reliable paths faster
- Start BFS from arbitrary node
- As soon as we came to already visited vertex
 - Align to alternative paths
 - Project low-covered path onto the main one


Velvet “tour bus” algorithm




Velvet “tour bus” algorithm




Velvet “tour bus” algorithm




Velvet “tour bus” algorithm




Velvet “tour bus” algorithm




Velvet “tour bus” algorithm




Velvet “tour bus” algorithm




Velvet “tour bus” algorithm




Velvet “tour bus” algorithm




Velvet “tour bus” algorithm




Velvet “tour bus” algorithm




Velvet “tour bus” algorithm



Velvet “tour bus” algorithm



Velvet “tour bus” algorithm



Velvet erroneous connection remover

- Erroneous connections don't have any recognized topological structure
- Have low coverage
- Removed using simple coverage cutoff


Homework

Implement de Bruijn graph


- Construction from FASTA/FASTQ
- Condensation
- Output to FASTA/DOT/...
- Simplification

Resolving repeats


Resolving repeats




Resolving repeats




Resolving repeats




Resolving repeats



Resolving repeats

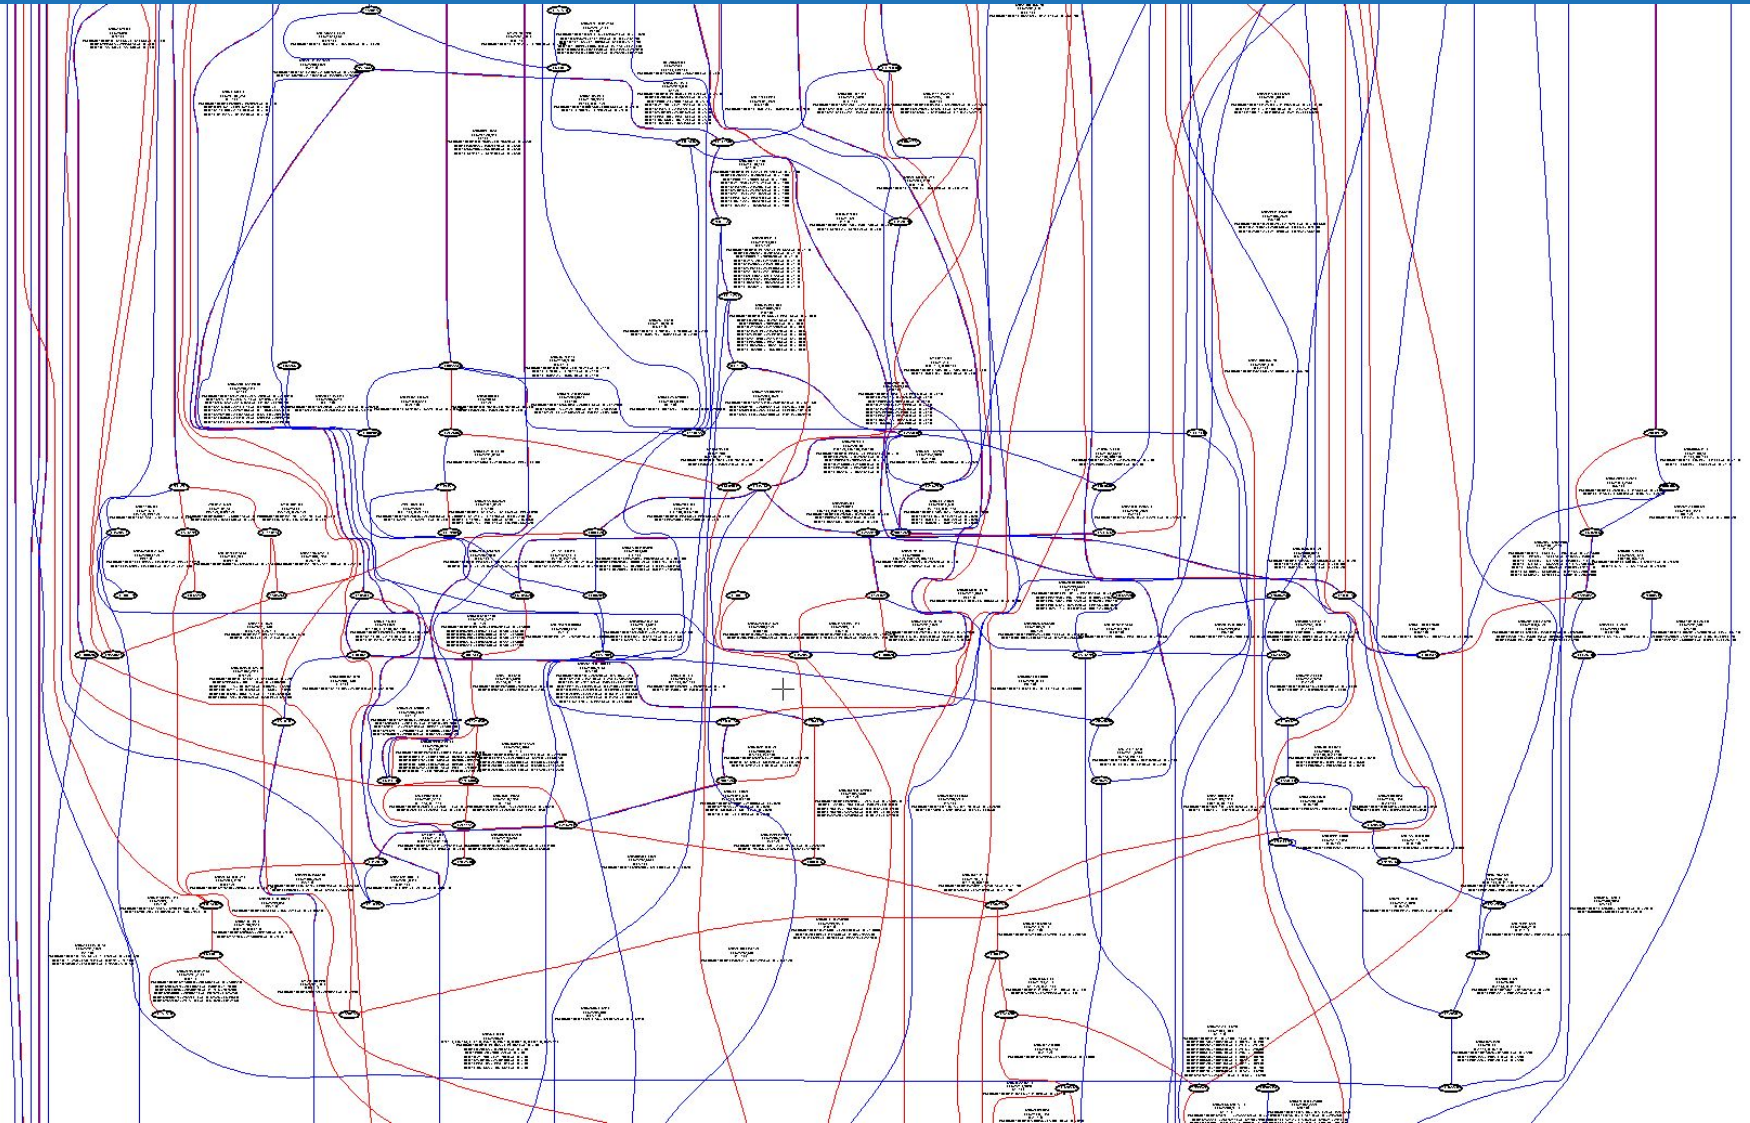


Resolving repeats




Real life

Part of *E.coli* genome, $K = 99$



Insert size distribution

Paired-end reads




Mate-pairs



Mate-pairs




3.5 kb DNA Molecule
with Biotin End Labels

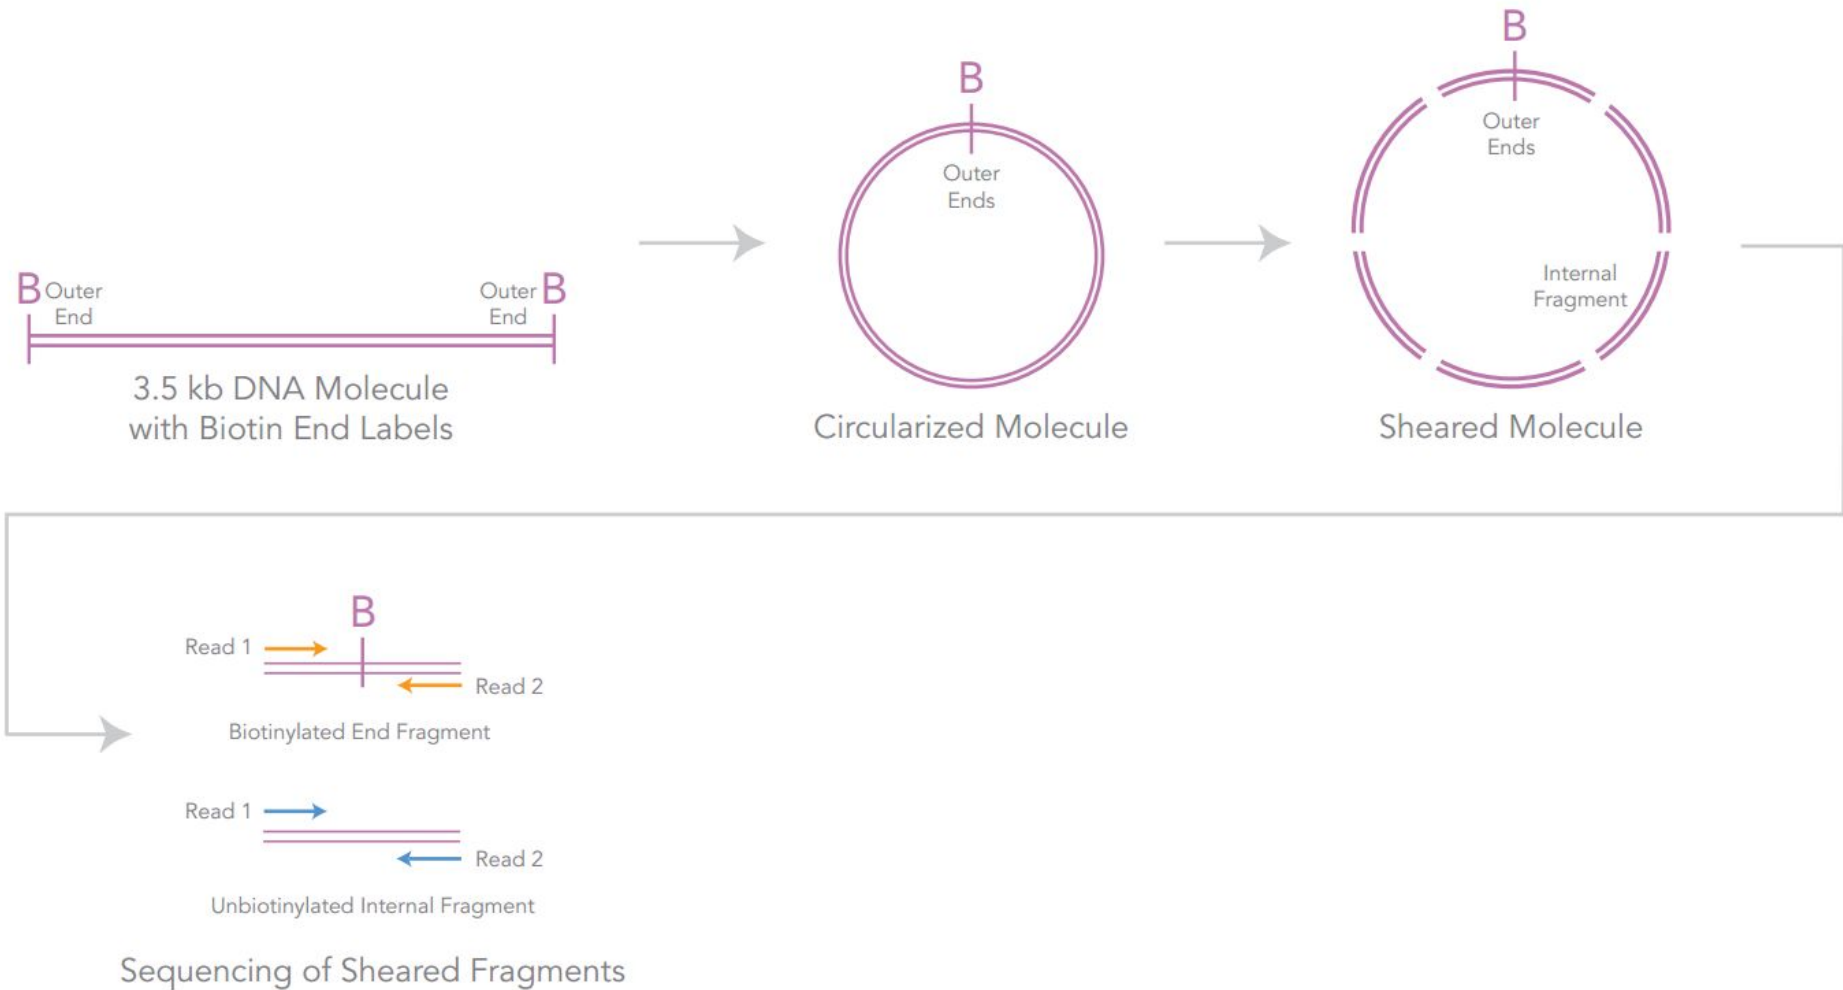


Circularized Molecule

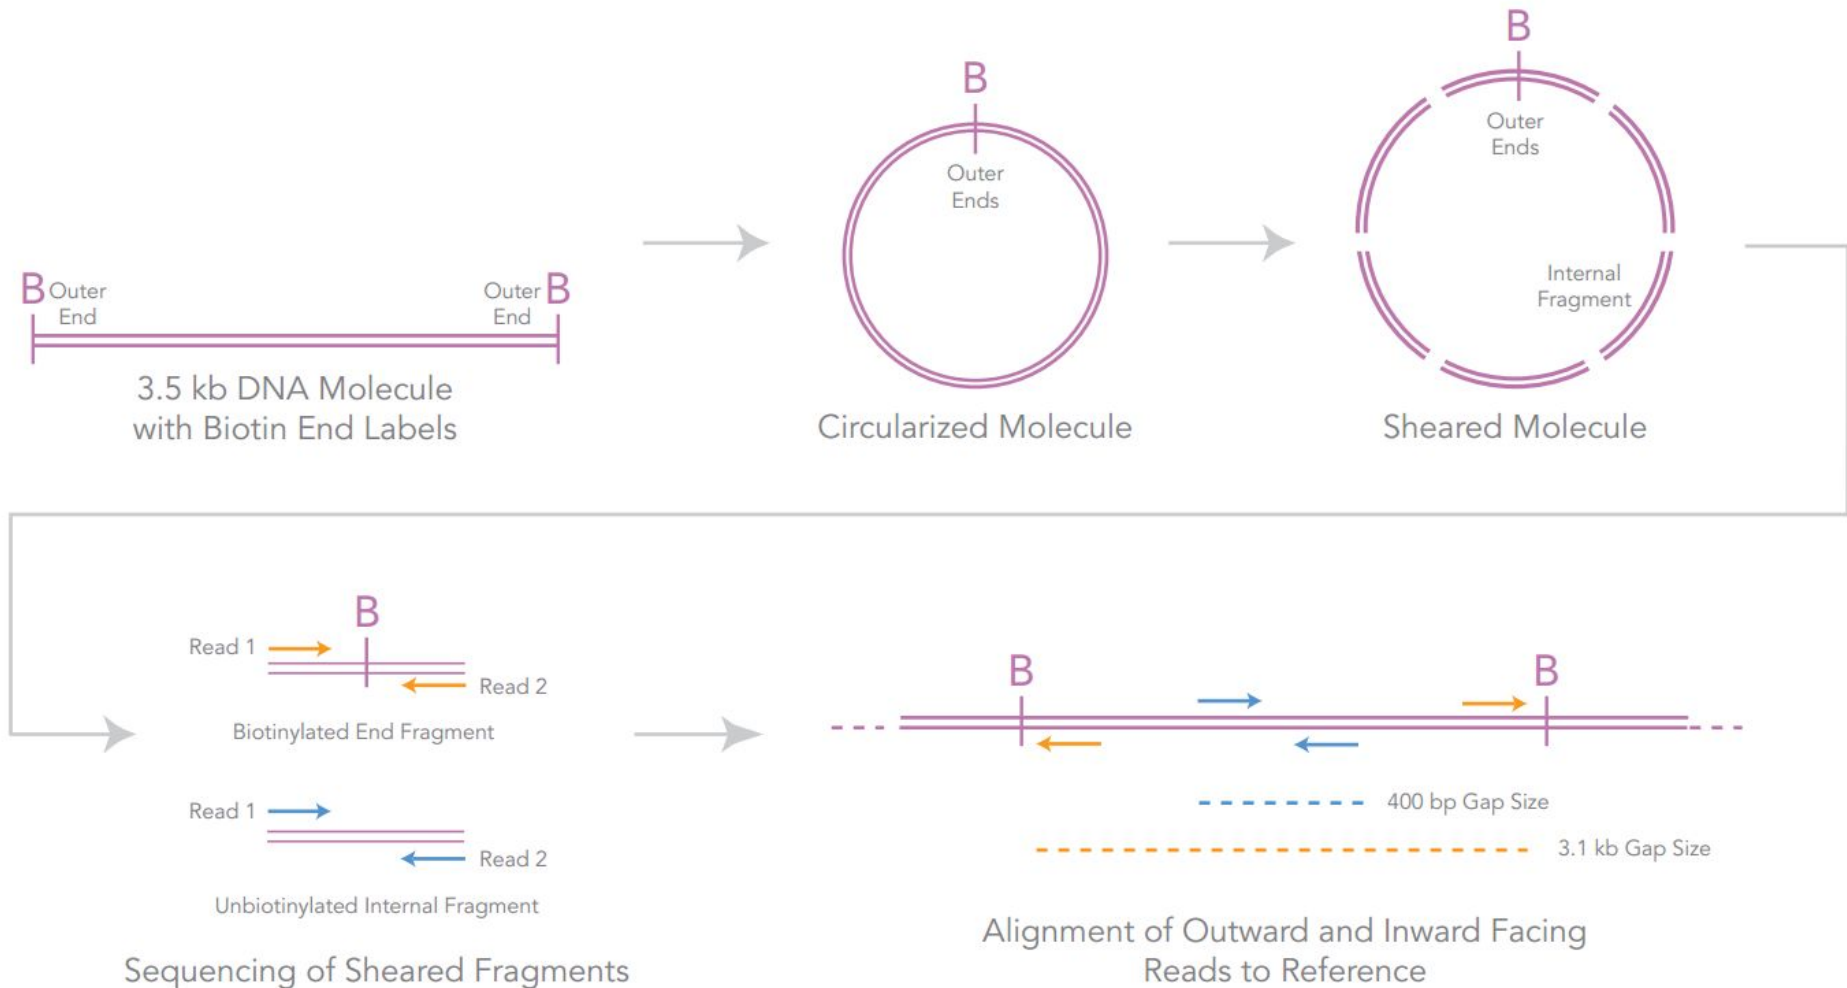
Mate-pairs



Mate-pairs




Mate-pairs




Mate-pairs

Conventional mate-pairs:




Mate-pairs


Conventional mate-pairs:




Illumina Nextera mate-pairs:




Long reads to the rescue




Long reads to the rescue









Long reads to the rescue



Long reads to the rescue



Sequencing technologies

	 SANGER SEQUENCING	 illumina®	 Roche 454 SEQUENCING	 ion torrent by life technologies™	 PACBIO®		 NANOPORE technologies™
Protocols		HiSeq, MiSeq			Subreads	CCS / HiFi	MinIon
Read length	500-900	25-300	400-1100	200-400	20K-100K	5K-20K	1K-3M
Error rate	0.001-0.1%	0.1-1%	1%	1-2%	2-10%	0.1-2%	5-15%
Error type	Indels & Mismatches	Mismatches only	Indels & Mismatches	Indels & Mismatches	Indels & Mismatches	Indels & Mismatches	Indels & Mismatches
Comments	Remains the golden standard	Error rate grows at the end of read	Problems with homopolymers	Problems with homopolymers	Errors distributed randomly	Error rate depends on sequencing settings	Typically several deletions in a row
Cost	\$\$\$\$\$	\$	\$	\$	\$	\$	\$

PacBio only assembly

Thm:

Perfect assembly possible iff

a) errors random

b) sampling is Poisson

c) reads long enough 2 solve repeats.

Note: e-rate not needed

Gene Meyers' twitter

New long reads vs Sanger assembly

High error rate => overlap detection is harder

- miniasm
- MHAP
- ...

Thank you!

Questions?